


EECS 376: Foundations of Computer Science

Lecture 07 – Greedy Algorithms



Greedy Algorithms

- A **greedy algorithm** is any algorithm that follows the problem-solving **heuristic** of making the locally optimal choice at each stage. ✱
- In many problems, a greedy strategy does not produce an optimal solution, but a greedy heuristic can yield locally optimal solutions that approximate a globally optimal solution in a reasonable amount of time.

Greedy Algorithms vs Dynamic Programming

- In contrast to dynamic programming, which carefully chooses a solution by considering the results of previous decisions, greedy algorithms do not look back. They make a series of choices that seem best at the moment, which can lead to suboptimal solutions for some problems.
- This is the main difference from dynamic programming, which is exhaustive and is guaranteed to find the solution.

Greedy Algorithms

Pick the best choice **NOW**.

Prove you end up with an optimal solution.



Proof Technique: Induction + “Exchange” argument


Warning: Greed is generally bad!



Greedy

Divide and conquer

Dynamic programming



- Fast
- Doesn't work for most problems

- Often slower than greedy and faster than DP
- Works when solutions to disjoint subproblems can be combined into final solution

- Generally slower (but still usually efficient)
- Applies to many problems

⋮

harder problems where none of these methods apply (coming soon)

Template

- Solve the problem in a “greedy”, “myopic” way
 - Rarely gives you an **exactly** optimum solution but makes for some very elegant algorithms when it does.
 - (Often works well for **approximation algorithms** — more on this later in the course.)
- **Main difficulty:** arguing for correctness
 - Exchange arguments

But sometimes greed can be good...

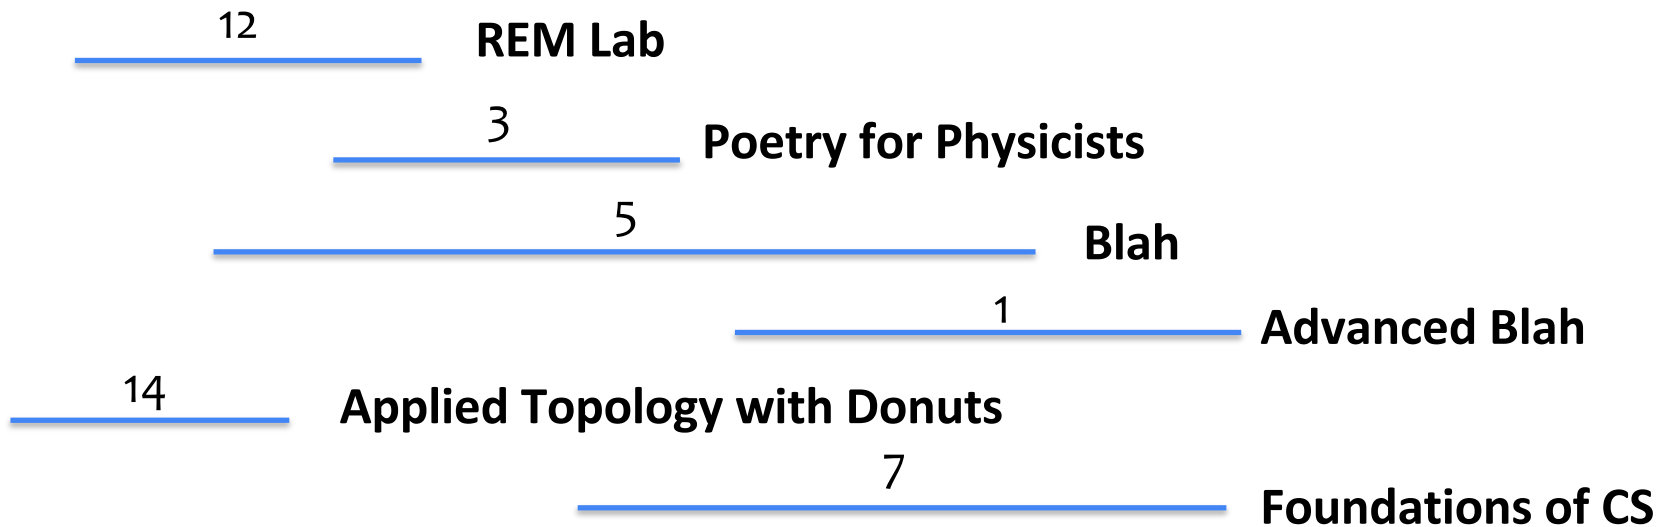
Unweighted Task Selection

Activity scheduling

- An **activity** i has **start time** s_i and **end time** f_i
- **Goal:** Given a set A of n activities (classes), select a subset $S \subseteq A$ that are **mutually disjoint** that maximizes $|S|$, i.e. a maximum **schedule**.
- Activities i and j are **disjoint** if their intervals $[s_i, f_i)$ and $[s_j, f_j)$ don't overlap
 - $s_i \geq f_j$ or $s_j \geq f_i$

Weighted Course Registration

(aka **Weighted** Task Selection)

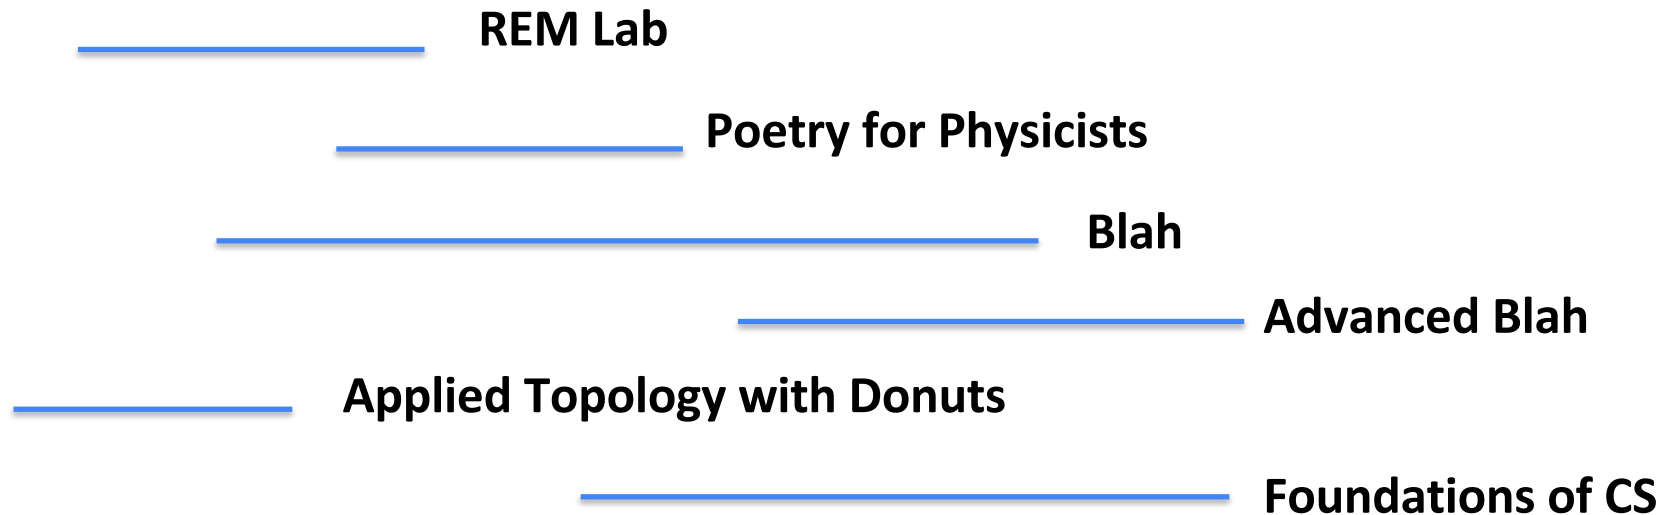


Goal: Choose a set of non-intersecting courses **with largest total value**.
(there may be many optimal solutions, we just seek one)

Recall this problem. We solved this using dynamic programming.
But in the **unweighted** case, we will show a simpler greedy algorithm.



Unweighted Course Registration

(aka Task Selection)




Goal: Choose a **largest** possible set of non-intersecting courses (there may be many optimal solutions, we just seek one!)

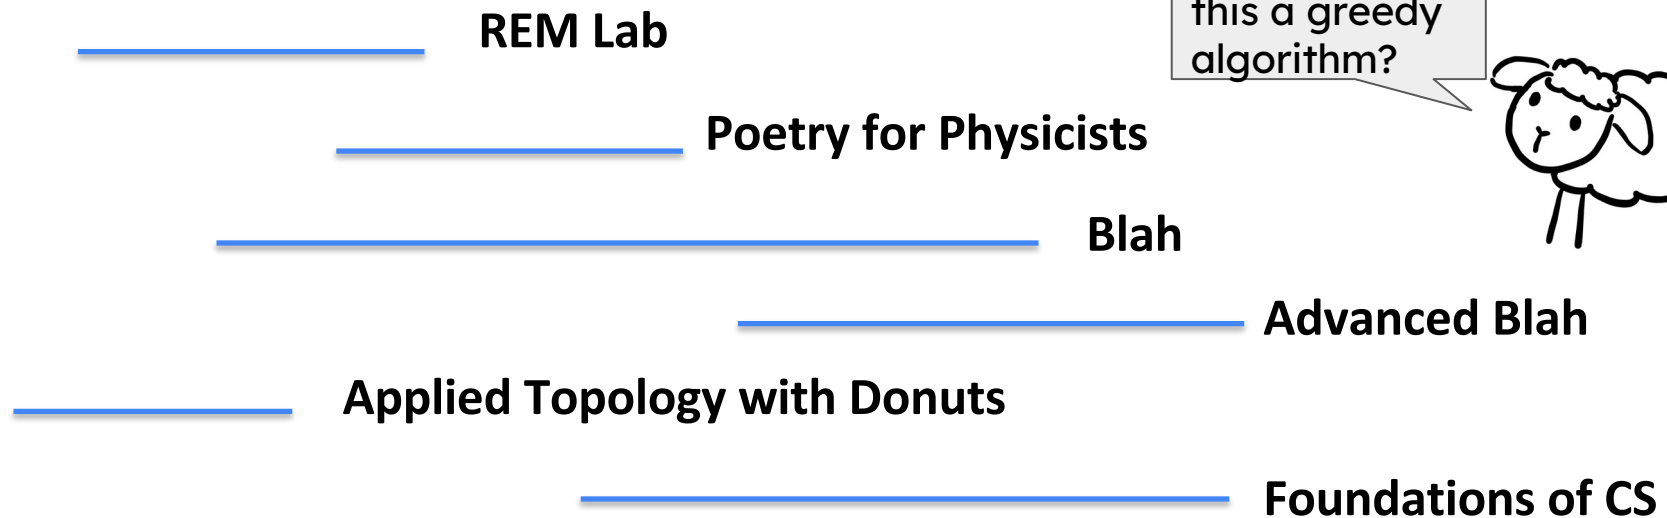
I have not 1, not 2, but 3
greedy algorithms!



Professor Y's Greedy Algorithms



X **Attempt 1:** Choose the **shortest interval** (breaking ties arbitrarily), take it, remove overlaps, recurse on remaining problem.




What makes this a greedy algorithm?




Counterexample:

5/16/24




Professor Y's Greedy Algorithms




✗ **Attempt 2:** Choose the **interval that starts earliest** (breaking ties arbitrarily), take it, remove overlaps, recurse on remaining problem.

Counterexample:




Professor Y's Greedy Algorithms



X Attempt 3: Choose the **interval that overlaps with the fewest other intervals** (breaking ties arbitrarily), take it, remove overlaps, recurse on remaining problem.

Counterexample:



No preeminent “Greedy” algorithm

- Possible greedy heuristics:
 - Pick a set of activities one at a time, *shortest activities first* (minimizing $|f_i - s_i|$).
 - Pick a set of activities one at a time, *earliest starting time* first.
 - Pick a set of activities one at a time, *earliest ending time* first.

A greedy algorithm: “Earliest Ending Time (EET)”


Assume they're sorted in increasing order by finishing time: $f_1 \leq f_2 \leq \dots \leq f_n$

```
Greedy( $s, f$ ):  
   $S \leftarrow \{1\}$   \\ chosen activities  
   $j \leftarrow 1$     \\ activity chosen with the largest  
   $f_j$   
  for  $i = 2..n$ :  
    if  $s_i \geq f_j$  :  
       $S \leftarrow S \cup \{i\}$   
       $j \leftarrow i$   
  return  $S$ 
```

Runtime: $O(n)$

A greedy algorithm: “Earliest Ending Time (EET)”

$t = 0$



```
Greedy( $s, f$ ):  
 $S \leftarrow \{1\}$   
 $j \leftarrow 1$   
for  $i = 2..n$ :  
    if  $s_i \geq f_j$ :  
         $S \leftarrow S \cup \{i\}$   
         $j \leftarrow i$   
return  $S$ 
```

A Correct Greedy Algorithm: “Earliest Ending Time (EET)”

- Sort the intervals by *ending time*
- Greedily take the **interval I that ends first** (break ties arbitrarily)
- Remove the intervals that overlap with the one just selected
- Recurse on remaining problem




A Correct Greedy Algorithm: “Earliest Ending Time (EET)”

- Sort the intervals by *ending time*
- Greedily take the **interval I that ends first** (break ties arbitrarily)
- Remove the intervals that overlap with the one just selected
- Recurse on remaining problem




A Correct Greedy Algorithm: “Earliest Ending Time (EET)”

- Sort the intervals by *ending time* / $n \log n$
- Greedily take the **interval I that ends first** (break ties arbitrarily)
- Remove the intervals that overlap with the one just selected
- Recurse on remaining problem



Let's see the big idea
of the proof of
correctness first and
then the formal
proof afterwards




A Correct Greedy Algorithm: “Earliest Ending Time (EET)”

Key Claim: The interval I that ends first is a **safe** choice i.e. it is in some optimal solution.

Why? Consider an optimal solution OPT . Let I_{OPT} be the interval that ends first in OPT .

- I_{OPT} ends at least as late as I .
- All other intervals in OPT start after I_{OPT} ends, and thus after I ends.
- Thus, we can take OPT and **exchange** I_{OPT} for I and get a valid solution of the same size!



Formal Proof of Correctness by Induction

intervals in order of addition (so, increasing end times)

Let I_1, I_2, I_3, \dots be the output of the EET algorithm


Goal: Prove that for all k , $I_1, I_2, I_3, \dots, I_k$ is in some optimal solution.

Proof by induction on k :

Base case: $k=0$.

Inductive hypothesis: Suppose $I_1, I_2, I_3, \dots, I_k$ is in some optimal solution $I_1, I_2, I_3, \dots, I_k, I'_{k+1}, I'_{k+2}, \dots$

Inductive step: Goal: show $I_1, I_2, I_3, \dots, I_{k+1}$ is in some optimal solution.



Formal Proof of Correctness by Induction


Inductive step: Goal: show $I_1, I_2, I_3, \dots, I_{k+1}$ is in some optimal solution.

By the inductive hypothesis, there exists an optimal solution:

$$\text{OPT} = I_1, I_2, I_3, \dots, I_k, I'_{k+1}, I'_{k+2}, \dots$$

Use an **exchange** argument as before!

- I'_{k+1} ends at least as late as I_{k+1} .
- All other intervals I'_{k+2}, \dots start after I'_{k+1} ends, and thus after I_{k+1} ends.
- Thus, we can take OPT and **exchange** I'_{k+1} for I_{k+1} and get a valid solution of the same size!



General Strategy commonly used for analyzing greedy algorithms:

Proof by induction using an “**exchange**” argument

The idea: Show that we can transform any **optimal solution** into the **solution given by our algorithm** by **exchanging** each piece of it out one-by-one without increasing the cost.

Key part of proof: Show that my greedy choice is **safe** i.e. it is in some optimal solution.


Induction just formalizes the idea that *each successive choice* is **safe**.

Minimum Spanning Trees

A Highway Problem

Input: an undirected graph with positive edge weights
e.g. a set of cities and distances between them

Output: minimum total length of highway to connect all cities
i.e. it should be possible to drive from any city to any other
using just the highways



A Highway Problem


Claim: Solution is a tree.

Why?

This fact will be useful later too

If a connected graph has a cycle, we can delete an arbitrary edge from the cycle and still have a connected graph.

Definition (review): A tree is a connected graph with no cycles.




The Highway Problem is commonly known as:

Minimum Spanning Tree (MST)

Given a graph G , a **spanning tree** is a subgraph of G that is spanning (uses all vertices) and is a tree.

MST problem: Given an undirected graph with positive edge weights, find a minimum weight spanning tree.



MST applications in computer science

Minimum spanning trees have a wide range of applications in computer science. Here are some notable examples:

- **Network Design:** Minimum spanning trees are crucial in the design of networks, such as telephone, electrical, and computer networks, to ensure minimal wiring with maximum connectivity
- **Routing Protocols:** In computer networking, minimum spanning trees are used to prevent loops in network routing. **The Spanning Tree Protocol (STP)** is a network protocol that ensures a loop-free topology for **Ethernet networks**.

MST applications...

Image Segmentation: In image processing, minimum spanning trees can be used to segment an image into different regions, which is useful for object recognition and other tasks.

Cluster Analysis: Minimum spanning trees can help in cluster analysis by connecting points into a tree structure based on their proximity, which can then be used to identify natural groupings within the data.

Approximation Algorithms: For NP-hard problems like the **traveling salesperson problem**, minimum spanning trees can be used to create approximation algorithms that provide near-optimal solutions.

MST applications...

Bioinformatics: In bioinformatics, spanning trees are used to construct phylogenetic trees, which represent the evolutionary relationships between different species.

Facility Location: Spanning trees can assist in determining the optimal locations for facilities like warehouses or power plants within a network.


Geographic Information Systems (GIS): They are used in GIS to create maps that minimize the total distance between locations.

Minimum Spanning Tree (MST)


Template for greedy algorithm:

Greedly pick an edge and add it to our MST. Repeat.

But which edge do we pick?



We need to pick a **safe** edge: an edge that appears in *some* optimal solution with all of the previously picked edges.



Kruskal's algorithm

- **Kruskal's algorithm** finds a minimum spanning forest of an undirected edge-weighted graph. If the graph is **connected**, it finds a **minimum spanning tree**. It is a **greedy** algorithm that adds to the forest the lowest-weight edge in each step that will not form a **cycle**.
- The algorithm's key steps are **sorting** and using a **disjoint-set data structure** to detect **cycles**. Its running time is dominated by the time to sort all of the graph edges by their weight.

Kruskal's algorithm (Cont'd)

- A **minimum spanning tree** of a **connected** weighted graph is a **connected** subgraph, without **cycles**, for which the sum of the weights of all the edges in the subgraph is **minimal**.
- For a **disconnected** graph, a **minimum-spanning forest** is composed of a **minimum-spanning tree** for each **connected** component.
- This algorithm was first published by **Joseph Kruskal** in 1956 and was rediscovered soon afterward by Loberman & Weinberger (1957).

Kruskal's Algorithm (1956)

Pick the minimum weight edge!



Kruskal(G): // G is a weighted, undirected graph

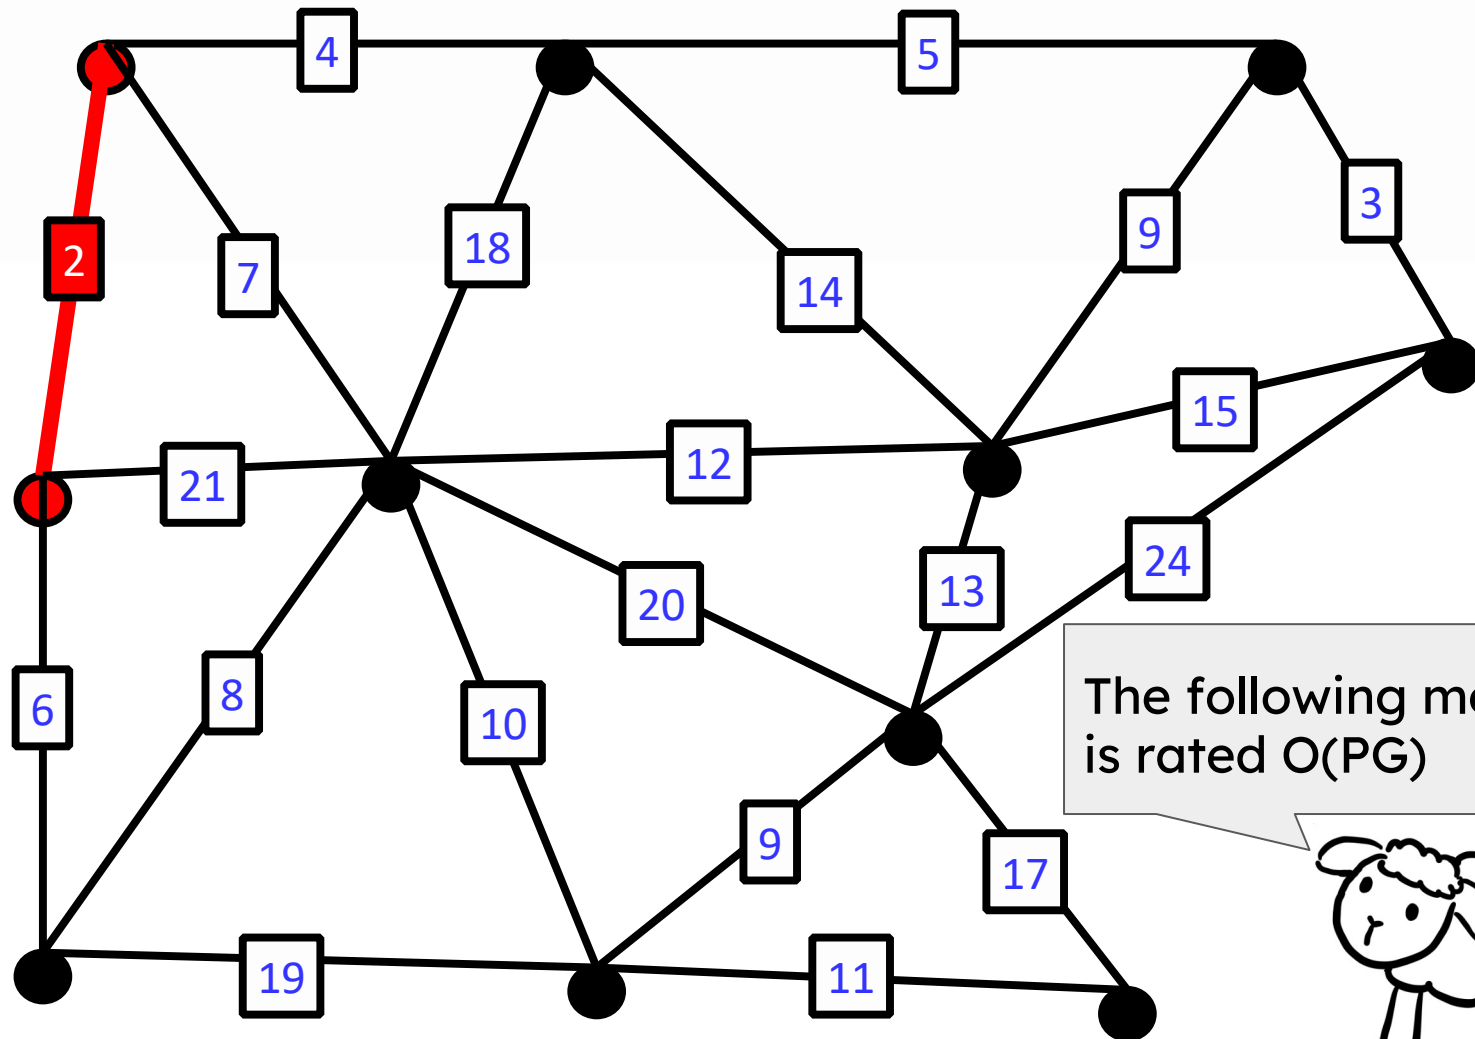
$T \leftarrow \emptyset$ // invariant: T has no cycles

for each edge e in increasing order of weight:


if $T + e$ is acyclic: $T \leftarrow T + e$

return T


Sorted weights: 2,3,4,5,6,7,8,9,9,10,11,12,13,14,15,17,18,19,20,21,24




Sorted weights: 2,**3**,4,5,6,7,8,9,9,10,11,12,13,14,15,17,18,19,20,21,24




Sorted weights: 2,3,4,5,6,7,8,9,9,10,11,12,13,14,15,17,18,19,20,21,24




Sorted weights: 2,3,4,5,6,7,8,9,9,10,11,12,13,14,15,17,18,19,20,21,24




Sorted weights: 2,3,4,5,6,7,8,9,9,10,11,12,13,14,15,17,18,19,20,21,24




Sorted weights: 2,3,4,5,6,7,8,9,9,10,11,12,13,14,15,17,18,19,20,21,24




Sorted weights: 2,3,4,5,6,7,8,9,9,10,11,12,13,14,15,17,18,19,20,21,24




Sorted weights: 2,3,4,5,6,7,8,9,9,10,11,12,13,14,15,17,18,19,20,21,24




Sorted weights: 2,3,4,5,6,7,8,9,**9**,10,11,12,13,14,15,17,18,19,20,21,24




Sorted weights: 2,3,4,5,6,7,8,9,9,**10**,11,12,13,14,15,17,18,19,20,21,24



Sorted weights: 2,3,4,5,6,7,8,9,9,10,**11**,12,13,14,15,17,18,19,20,21,24



Sorted weights: 2,3,4,5,6,7,8,9,9,10,11,**12**,13,14,15,17,18,19,20,21,24



The complexity of Kruskal's algorithm

- **Time Complexity:** The time complexity of **Kruskal's algorithm** is $O(E \log V)$ where E is the number of edges and V is the number of vertices in the graph. This complexity arises because the algorithm needs to sort all the edges of the graph, which takes $O(E \log V)$ time, and then perform union-find operations on the edge set.
- **Space Complexity:** The space complexity of **Kruskal's algorithm** is $O(V + E)$. This accounts for the space needed to store the graph's edges and vertices.

Kruskal's Algorithm: Correctness

Why does Kruskal's algorithm return a **spanning tree**?

1. Why is it a **tree**?
2. Why is it **spanning** (i.e. connects all vertices)?

Now we need to show it returns a **minimum** spanning tree.

Let's assume the weights are distinct by breaking tie arbitrarily.

Proof of Minimality: Warm up

Let T be a spanning tree.

Let e be an edge not in T



Q: How does $T + e$ look like?

A: a cycle C + a forest (i.e., many trees)

Let e' be an edge in C

Q: How does $T + e - e'$ look like?

A: Another spanning tree T'



(前提: weights are distinct)


Proof of Minimality: Key Lemma

Key Lemma:

- For any set $S \neq V$ of vertices.
- the min-weight edge e crossing S must be in MST of G .

Proof:

- Let T be an MST of G .
- Suppose $e \notin T$ for contradiction.
- T must cross S .
- Consider the cycle C in $T + e$.
- There is another edge $e' \in C \cap T$ crossing S .
- Consider a spanning tree $T' = T + e - e'$.
- As $w(e) < w(e')$, T is not MST. Contradiction. QED



(前提: weights are distinct)

Proof of Minimality: Finish

Thm: The output tree T of Kruskal is an MST.

Proof:

- For each edge e chosen by Kruskal,
- Can you find a set S where e is a min-weight edge crossing S ?
 - Yes. How?
 - Let $e = (u, v)$. Let T_u be the tree that correctly contains u .
 - $S =$ vertex set of T_u .
- So, every edge $e \in T$ must be in MST.
- So T is minimum.

It's proof

Removing the distinct-weight assumption

- We assumed distinct edge weights. Why is this valid?
- If a graph G have edges with same weight,
 - Just break tie arbitrarily, say by the lexicographical order.
- The weight of the minimum spanning tree does not change at all.

Setting up the Induction

edges in order of addition (so, **non-decreasing** weights)

Let $T = e_1, e_2, e_3, \dots$ be the output of Kruskal's algorithm

Goal: Prove that for all k , the edge set $e_1, e_2, e_3, \dots, e_k$ is in some MST.

Proof by induction on k :

Base case: $k=0$.

Inductive hypothesis: Suppose the edge set $e_1, e_2, e_3, \dots, e_k$ is in some MST $T' = e_1, e_2, e_3, \dots, e_k, f_1, f_2, \dots$ ← f edges listed in no particular order

Inductive step: Goal: Show the edge set $e_1, e_2, e_3, \dots, e_{k+1}$ is in some MST.

Inductive step


Inductive step: Goal: Show the edge set $e_1, e_2, e_3, \dots, e_{k+1}$ is in some MST.

By the inductive hypothesis, there exists an MST:

$$T' = e_1, e_2, e_3, \dots, e_k, f_1, f_2, \dots$$

Case 1: $e_{k+1} \in T'$

We are done.

根据 tree 定义, 把 e_{k+1} 加入 T' 中 一定会
导致一个 cycle 

Case 2: $e_{k+1} \notin T'$

Claim: We can perform an edge exchange: We can take T' , add e_{k+1} , and remove an edge in f_1, f_2, \dots to get a spanning tree that still has minimum weight.

且由于 T' 中有 e_{k+1} 但却没有这个 cycle, 

说明这个 cycle 中至少有一边是 T' 中没有的, 


于是取 T' 的此 cycle 中任意一个 T' 中没有的边 f , 用来替换 e_{k+1}

Proof of Claim:

Which edge f_i do we exchange with e_{k+1} ?

加入 e_{k+1} 在 T' 中形成的 cycle 中, 取任意一个 $f \in T$

After the exchange, why is the graph still connected?



After the exchange, why is the graph still a tree?



⇒ 替换后仍是 spanning tree

⇒ 由于 T 没考虑 f 却考虑了 e_{k+1} ,

$$w(f) \geq w(e_{k+1})$$

因而替换后, $\sum w$ 减小了 ⇒ T' 并不是 MST

⇒ contradicts ⁵⁸

因而不可能 ⇒ $e_{k+1} \in T'$ ✗

Question

Suppose that G has distinct edge weights.

1. MST(G) is unique. Why? 因为 $T = T'$ by induction
2. Let G' be obtained by doubling the weight of G .
Is it always the case that $MST(G) = MST(G')$?

Running time (we won't fully prove) of Kruskal's algorithm: $O(m \log n)$
⇒ need to analyze a data structure for detecting cycles (disjoint-sets data structure) (aka union-find data structure)

Best-known running time (Chazelle 2000): $O(m \cdot \alpha(m, n))$

Open problem: $O(m)$?

↑
Inverse Ackermann
function: an extremely
slow growing function:
 $\alpha(n) \leq 4$ even when n is
particles in known
universe

Seth Pettie and Vijaya Ramachandran (2002) gave an asymptotically **optimal** algorithm. But nobody knows how fast it is!



When does Greedy work?

- Although Greedy often does not work...
- There are **classes of problems** that Greedy will work
 - **Matroid:**
 - Greedy will give an optimal solution
 - Matroid captures a minimum spanning tree.
 - **Submodular maximization:**
 - Greedy will give an approximately optimal solution
- You can look up what they are.