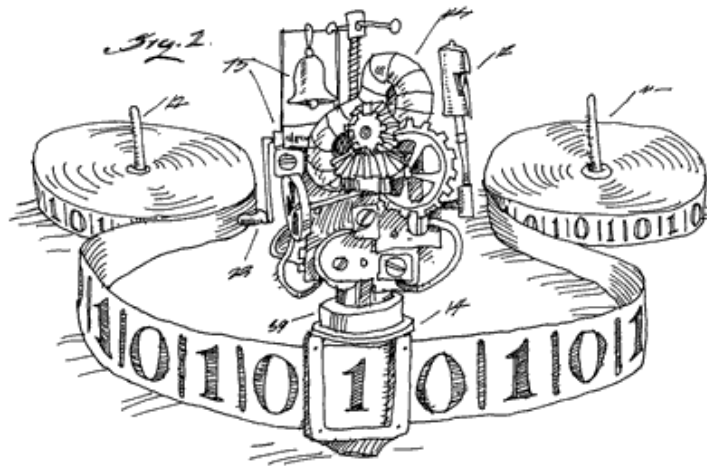


EECS 376: Foundations of Computer Science

Lecture 25 – Interactive Proofs & Zero-Knowledge proofs



Seminal Papers

- S. Goldwasser, S. Micali, C. Rackoff, **The knowledge complexity of interactive proof systems**, STOC '85: Proceedings of the Seventeenth Annual ACM Symposium on Theory of Computing, December 1985.
- S. Goldwasser, S. Micali, A. Wigderson, **Proofs that yield nothing but their validity and a methodology of cryptographic protocol design**, 27th Annual Symposium on Foundations of Computer Science, October 1986.

co-originators of the important concepts of "interactive proofs" and "zero-knowledge proofs"

Goldwasser and Micali win Turing Award

Team honored for 'revolutionizing the science of cryptography.'

Abby Abazorius, CSAIL
March 13, 2013

Including introducing
interactive proof systems and
zero-knowledge proofs



Chris Peikert's PhD advisor



Shafi Goldwasser

Silvio Micali

co-originators of the important concepts of "interactive proofs" and "zero-knowledge proofs"

C. Rackoff

Co-won the 1993 Godel Prize



A. Wigderson

Won the Turing Award last year!

For “reshaping our understanding of the role of randomness in computation”



Interactive Proof System

An **interactive proof system** is a concept in **computational complexity** theory that models computation as a dialogue between two parties: a **prover** and a **verifier**.

- **Prover**: This party has unlimited computational resources but cannot be **trusted**.
- **Verifier**: This party has limited computational power but is assumed to be always **honest**.

Interactive Proof System

- The **prover** and **verifier** exchange messages to determine whether a given **string** belongs to a **language**.
- The **interaction** continues until the verifier is **convinced** of the answer to the problem.

Properties of Interactive Proof Systems

Two main properties characterize **interactive proof systems**:

- **Completeness**: If the statement is true, an honest prover can convince the honest verifier of its truth.
- **Soundness**: If the statement is false, no prover can convince the honest verifier that it is true, except with some small probability.

Applications of Interactive Proof Systems

Interactive proof systems have a wide range of applications across various fields. These applications demonstrate their versatility in enhancing security, verifying correctness, and solving complex computational problems.

- **Cryptographic Protocols:** They are fundamental in designing protocols where trust and verification are crucial, such as in zero-knowledge.
- **Approximation Algorithms:** Interactive proofs help understand the power and limitations of approximation algorithms.

Applications of Interactive Proof Systems

- **Program Checking**: They are used to check the correctness of programs in a way that the verifier can be convinced of the program's correctness without having to execute the program.
- **Complexity Theory**: Interactive proofs have provided insights into the hardness of certain computational problems, such as graph isomorphism and the approximate shortest lattice vector problem.
- **Interactive Theorem Provers**: These are tools used to certify mathematical theories and construct machine-verified proofs.

Interactive Proof System

(with k -round interactions)

- Let $L \subseteq \{0,1\}^*$ be a language and (A, B) an **interactive pair of Turing machines**. We say that (A, B) is an *interactive proof system* for L if A (the **prover**) has **infinite power**, B (the **verifier**) is **polynomial time** and they satisfy the following properties.
 - For any $x \in L$ given as input to (A, B) , B halts and accepts with probability $1 - \frac{1}{n^k}$ for each k and sufficiently large n .
 - For any ITM A^* and any x not in L given as input to (A^*, B) , B accepts with probability at most $\frac{1}{n^k}$ for each k and sufficiently large n .

Interactive Complexity Classes

- We define **IP**, *Interactive Polynomial-time*, to be the class of languages possessing an **interactive proof system**.
- In this case, we may also say that L is interactively provable. To emphasize that the prover has unlimited power, we may write \mathbf{IP}_∞ for **IP**.

Graph Isomorphism

- We say two graphs G_1 and G_2 are **isomorphic** if they are the same up to a renumbering of vertices. In other words, if there is a permutation π of the labels of the nodes of G_1 such that $\pi(G_1) = G_2$
- The **graph isomorphism** problem is important in a variety of fields and has a **rich** history.
- Along with the **factoring problem**, it is the most famous NP problem that is not known to be either in **P** or **NP-complete**.

Protocol: Private-coin Graph (Non-)isomorphism

- Both **graph isomorphism** and its **complement** can be shown to belong to **IP** as follows.
- The protocol has k - round interactions. Each round proceeds as follows:
 - V: pick $i \in \{1, 2\}$ uniformly randomly. Randomly permute the vertices of G_i to get a new graph H . Send H to P.
 - P: identify which of G_1 and G_2 was used to produce H . Let G_j be that graph. Send j to V. V: accept if $i = j$; reject otherwise.

Interactive Complexity Classes

- If we allow the **probabilistic verifier** machine and the **all-powerful prover** to interact for a **polynomial** number of rounds, we get the class of problems called **IP**.
- In 1992, **Adi Shamir** was able to prove one of the **central** results of complexity theory that **IP** equals **PSPACE**, the class of problems solvable by an ordinary deterministic Turing machine in **polynomial space**.

Zero Knowledge



What is a Zero-Knowledge Proof?

A convincing demonstration that some statement is true—
without revealing anything beyond the fact that it's true!

Example: Suppose you knew that these two pens are different.
How could you convince me they're different without telling me
what is different about them?

Zero-Knowledge Proof for “Where’s Waldo?” (or “where’s the puffin among the penguins?”)

How could you convince me that there is a puffin in this picture without revealing its location (or anything else)?




Computer Scientist Explains One Concept in 5 Levels of Difficulty | WIRED

How are Zero-Knowledge Proofs Useful?

- **Electronic money / Blockchain:** Prove that your account has enough money for a transaction, without revealing your balance.
- **Group signatures:** Prove that a member of a certain group signed a message, without revealing which one did.
 - E.g., give students keycard access to restricted areas without tracking individual students.
- **Multi-party computation:** Compute aggregate functions of private data (e.g., medical records), without revealing individual data.

The Model

Both parties know x and L



Let's put aside zero knowledge for a moment and recall **verifying a certificate for an instance x of a language L in NP.**

Given an instance x
(of an NP-language L),
I send a certificate c

I run a verification
algorithm $V(x, c)$

c



Merlin
“the prover”
All-powerful but
untrustworthy


Recall that a language L is in NP if:

- $x \in L \Rightarrow$ Merlin can give Arthur a convincing “proof” that $x \in L$, i.e., exists certificate c so that $V(x, c)$ accepts
- $x \notin L \Rightarrow$ Merlin can’t “fool” Arthur into accepting that $x \in L$, i.e., for any c , $V(x, c)$ rejects.




Arthur
“the verifier”
Restricted to
poly-time
computation

Hamiltonian Cycle



I send a
Hamiltonian
Cycle in G



I check it, and am
convinced that G
has a Hamiltonian
Cycle



This is **NOT** a ZK proof because
Arthur learns an actual Ham Cycle
(not just the fact that G has a Ham Cycle)!

Merlin
“the prover”
All-powerful but
untrustworthy

Can we make this into a ZK proof?

YES!

By using randomness and interaction.

Arthur
“the verifier”
Restricted to
poly-time
computation

*Actually it's possible to make it non-interactive (we won't show)

Requirements for a protocol to be a ZK proof

1. **“Completeness”**: If $x \in L$, Merlin can cause Arthur to accept.
2. **“Soundness”**: If $x \notin L$, then Merlin cannot “fool” Arthur into accepting, except with some *small probability*. (← new, needed weakening)
3. **“Efficiency”**: Arthur runs in polynomial time (in the size of x).
4. **“Zero knowledge”**: Arthur “learns nothing” but the fact that $x \in L$.

Q: But what exactly does that mean? How to define it?

A: Given that $x \in L$, whatever Arthur saw from Merlin could have been generated by Arthur on his own, without ever interacting with Merlin.

↳ If Merlin uses randomness, Arthur could have generated messages from the same probability distribution as Merlin’s messages.

A Zero-Knowledge Proof for Hamiltonian Cycle

Merlin convinces Arthur that G has a Ham Cycle
without revealing anything else



A Zero-Knowledge Proof for HamCycle [Blum '86]

Manuel Blum, How to Prove a Theorem So No One Else Can Claim It, Proceedings of the International Congress of Mathematicians, Berkeley, California, USA, 1986.




Some academic descendants of Blum in our department:
Chris Peikert, Euiwoong Lee, Ben Fish, Wei Hu

The Protocol has k - round interactions. Each round proceeds as follows



I am supposed to:

- Randomly shuffle vertices.
- Write adjacency matrix.
- Cover each entry with a post-it note, and send!



I randomly choose one *challenge*:

- “**reveal all**” or
- “**reveal cycle**”

“**reveal all**”


If “**reveal all**”, I reveal the *entire* adjacency matrix *and* the shuffling permutation

	1	2	3	4	5	6
1	0	1	1	1	0	0
2	1	0	0	0	1	0
3	1	0	0	0	1	1
4	1	0	0	0	1	1
5	0	1	1	1	0	1
6	0	0	1	1	1	0

permutation: 4, 6, 3, 5, 2, 1

Accept if this matches original graph G

The Protocol has k - round interactions. Each round proceeds as follows



Quiz:


1. How would Merlin fool Arthur if **“reveal cycle”**?
2. How about **“reveal all”**?
3. What if Merlin didn't shuffle vertices?

I randomly choose one *challenge*:

- **“reveal all”** or
- **“reveal cycle”**

I am supposed to:

- Randomly shuffle vertices.
- Write adjacency matrix.
- Cover each entry with a post-it note, and send!




“reveal cycle”



Accept if revealed entries are 1s, and correspond to a Ham Cycle (ignoring G itself)²⁵

If **“reveal cycle”**,
I reveal only the
(shuffled) Ham
Cycle




Analysis

“Efficiency”: Arthur runs in polynomial time

- If “**reveal cycle**”, check in poly-time that the revealed part is a HamCycle

	1				
				1	
1					
					1
			1		
		1			

- If “**reveal all**”, check in poly-time that the initial graph and the revealed graph are the same using the permutation




Arthur had this
from the beginning

permutation: 4, 6, 3, 5, 2, 1

	1	2	3	4	5	6
1	0	1	1	1	0	0
2	1	0	0	0	1	0
3	1	0	0	0	1	1
4	1	0	0	0	1	1
5	0	1	1	1	0	1
6	0	0	1	1	1	0

Arthur is then given this



Just relabel and check

Analysis

“Completeness”:

If G has a Ham Cycle, **Merlin** can cause **Arthur** to accept:

Merlin follows the protocol.

If “**reveal cycle**”, **Arthur** can check the existence of HamCycle and accept.

If “**reveal all**”, **Arthur** can check that the two graphs are the same and accept.

“Soundness”:

If G has **no** Ham Cycle, **Merlin** “fools” **Arthur** into accepting with $\text{prob} \leq \frac{1}{2}$:

If **Merlin modified** the graph, **Merlin** fails if “**reveal all**”.

If **Merlin did not modify** the graph, **Merlin** fails if “**reveal cycle**”

In both cases, **Arthur** is fooled with $\text{prob} \leq \frac{1}{2}$

How can we reduce this “fooling” probability to near 0?


Analysis

“Zero knowledge”:

If G has a Ham Cycle,
then **Arthur**, *without ever interacting with Merlin*, could have generated messages from the same probability distribution as **Merlin's**.

If **“reveal cycle”**: Arthur sees a HamCycle with randomly permuted vertices.

- Can he generate this on his own?
- Yes!
- **Note:** Arthur does not need to find HamCycle in the original graph




If **“reveal all”**, Arthur sees a random permutation and the original graph after vertex relabeling

- Can he generate this on his own?
- Yes!

permutation: 4, 6, 3, 5, 2, 1

	1	2	3	4	5	6
1	0	1	1	1	0	0
2	1	0	0	0	1	0
3	1	0	0	0	1	1
4	1	0	0	0	1	1
5	0	1	1	1	0	1
6	0	0	1	1	1	0



28

Analysis

- Can you implement this protocol on the internet?
 - Merlin and Arthur are far away from each other.
 - What is missing?
- Not clear how to implement the digital version of “post-it”!

Digital Commitment Schemes

A digital version of this ZK proof requires a “**commitment scheme**”.

I.e., **Merlin** “commits” to bits without revealing them to **Arthur**, and can later reveal any of them as needed—without being able to *change* them.

This is possible!
(under cryptographic assumptions)

We won't prove it, but here's some intuition:

- **Merlin** “commits” to primes p, q by sending **Arthur** their product $n=pq$.
- **Arthur** can't factor efficiently, so doesn't know p and q .
- **Merlin** can later reveal p, q . **Arthur** checks that they are prime and that their product is n . No other numbers satisfy these properties, so Merlin could not have changed them!

(Needs to be extended to allow **Merlin** to commit to a desired *bit*.)

Zero-Knowledge Proofs for NP (and Beyond)

There is a ZK proof for ANY problem in NP!
(under cryptographic assumptions)

How?

Let L be a language in NP.
We want a ZK proof that $x \in L$.

Let f be a p.t.m.reduction from L to HC
(exists since HC is NP-complete).

So, $x \in L \Leftrightarrow f(x) \in \text{HC}$.

This includes NP-complete problems, problems in P, decision versions of discrete log and factoring, ...



To ZK-prove that $x \in L$, use the ZK proof for HC to prove that $f(x) \in \text{HC}$.

First ZK proof for an NP-C problem (with larger error) [Goldwasser-Micali-Wigderson '86]

[Goldwasser-Micali-Rackoff '85]

Beyond NP

Theorem (we won't show): Under cryptographic assumptions (in particular, the existence of OWFs), the set of problems with a **ZK proof** is *equivalent* to those in **PSPACE**, i.e., solvable w/ **polynomial space**.

[Lund-Karloff-Fortnow-Nisan-Shamir'92]

It is known that **NP** \subseteq **PSPACE**, and it is conjectured that **NP** \neq **PSPACE**.

An Unconditional Zero-Knowledge Proof

Unconditional ZK-Proof

Two graphs are *isomorphic* if one can be obtained from the other by relabeling the vertices.

Problem: Graph Non-Isomorphism


Given two graphs, are they “different” (non-isomorphic)?

Theorem:

There is a ZK proof for **Graph Non-Isomorphism** [GMW '86]

A ZK Proof for Graph Non-Isomorphism

The Protocol has k rounds. Each round proceeds as follows



I say whether
it's a relabeled
of G1 or G2

	1	2	3	4	5	6
1	0	1	1	1	0	0
2	1	0	0	0	1	0
3	1	0	0	0	1	1
4	1	0	0	0	1	1
5	0	1	1	1	0	1
6	0	0	1	1	1	0

(a scrambling
of G1, generated
by Arthur)

1. I randomly choose G1 or G2,
2. randomly relabel the vertices
3. send the resulting adjacency matrix!

It's G1!



Quiz: What if Arthur did not randomly relabel vertices?
How would Merlin fool him?

This is reminiscent
of the pen example



I accept if
Merlin got the
right answer

Analysis

Merlin can convince Arthur that G_1 and G_2 are non-isomorphic without revealing anything else.

“Completeness”:

If G_1 and G_2 are **non-isomorphic**, Merlin can cause Arthur to accept:

Merlin can tell which relabeled graph he got,
Arthur see a correct answer.

“Soundness”:

If G_1 and G_2 are **isomorphic**, Merlin “fools” Arthur into accepting with prob $\leq \frac{1}{2}$:

Merlin can't tell which scrambled graph he got,
Merlin's answer is correct with probability $\frac{1}{2}$.

How can we reduce this “fooling” probability to near 0?

Analysis

Merlin can convince Arthur that G_1 and G_2 are non-isomorphic without revealing anything else.

“Zero knowledge”:

If G_1 and G_2 are non-isomorphic, then whatever **Arthur** saw from **Merlin** could have been **generated by Arthur on his own, without ever interacting with Merlin.**

Arthur already knows the correct graph.
He can generate the correct answer.

“Efficiency”: **Arthur** runs in polynomial time

Trivial.

Wrap Up

Reflection on “Zero knowledge”

Informal:

Arthur “learns nothing” but the fact that $x \in L$.

Formal:

Given that $x \in L$,
whatever **Arthur** saw from **Merlin** could have been
generated by **Arthur** on his own,
*without ever interacting with **Merlin**.*

We can formalize vague concepts, like knowledge,
via **computational power**.

This is a deep insight from theoretical CS.

More example: What is **pseudo-random**? What is **private**? What is **secure**?
See [Imitation Games - Avi Wigderson](#)

Randomized Verifier is Essential

Claim: If a language L has a **ZK proof** where the verifier is **deterministic**, then L is in **P**.

Bottom line: if the verifier isn't randomized, then the prover cannot prove anything *in a zero-knowledge manner* than what the verifier can already solve on his own. So, it is not useful at all.

Randomized Verifier is Essential

Claim: If a language L has a **ZK proof** where the verifier is **deterministic**, then L is in P .

Without loss of generality, we can assume this:

1. There is only a single one-way message from Merlin.
 - Merlin knows Arthur's (deterministic) responses anyway.
 - So just concatenate all of Merlin's messages into one.
2. The strong **"Soundness"**:
If $x \notin L$, then Merlin "fool" Arthur with probability 0.
 - Otherwise, if fooling prob > 0 , there is a way to fool Arthur prob 1 (as Arthur is deterministic)
3. Merlin is deterministic.

So, w.l.o.g., Arthur is exactly the Verifier in the setting of NP problem

Randomized Verifier is Essential

Claim: If a language L has a **ZK proof** where the verifier is **deterministic**, then L is in **P**.

Since Arthur is exactly the Verifier in the setting of NP problem

- If x is in L , Merlin can send a message C and convince Arthur to accept.
- If x is not in L , for any Merlin's message C , Arthur always rejects.
- By the zero-knowledge condition, Arthur can come up with C in poly time without Merlin's help.
- So, Arthur can solve the problem in polynomial time on his own! L is in **P**.

Admin

Exam Review Session led by Erik:
Tuesday, June 25th, 1-3 pm, BBB 1670.

HW5 is due today at 8 pm.

Reminder: Filling out the course evaluations is 1% of your grade, which is otherwise covered by the final exam. Remember to submit the receipt on Gradescope.