

Dynamic Programming for Shortest Paths in Graphs


Did somebody
say G-raph?




Why Graphs?




The Internet



Social Networks



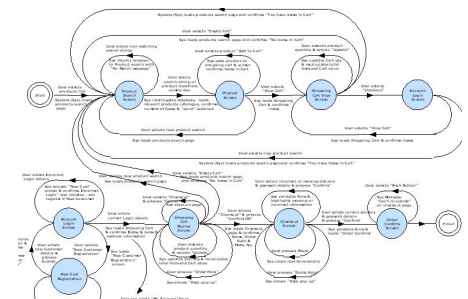
Biological and Chemical Networks



Circuits




Transportation Networks




State Transition Networks

Directed and undirected graphs

Directed graph




Undirected graph



Distance from s to t (denoted $\text{dist}(s,t)$) = minimum over all paths P from s to t , of the sum of the edge weights along P .

Notation: V = vertex set, E = edge set, $n = |V|$, $m = |E|$,

Why do we even care about negative weights?




The shortest paths problems we'll consider

Input: Directed weighted graph. Weights can be negative but assume no negative-weight cycles.


Single-Source Shortest Paths (SSSP): Given a special vertex s , find a shortest path from s to every vertex t .

All-Pairs Shortest Paths (APSP): For every pair s, t of vertices, find a shortest path from s to t .



	1	2	3	4	5
1	0	1	2	0	∞
2	0	0	2	-1	∞
3	-2	-1	0	-2	∞
4	1	2	3	0	∞
5	6	6	8	5	0


What about
single-pair
shortest path?




Two Key Observations

Principle of
Optimality


1. If a shortest path from s to t goes through vertex v , then it must take a **shortest path from s to v** , then a **shortest path from v to t** .
2. Since there are no negative-weight cycles in the graph, there is a shortest path from s to t with **no cycles** in it.



Check out my recurrence for SSSP!




Your equation is technically correct, but it's not really a recurrence and it doesn't work for DP.



$$\text{dist}(s,v) = \min_{(u,v) \in E} \{ \text{dist}(s,u) + \ell(u,v) \}$$

In the shortest $s \rightarrow v$ path,
 u is the last vertex before v
(and u could be s)




Base case: $\text{dist}(s,s) = 0$

Notation:

- $\ell(y,z)$ is the weight (or “length”) of the edge $y \rightarrow z$
- $\text{dist}(y,z)$ is the distance from y to z

Recurrence for SSSP

Key Definition. Let $\text{dist}^{(i)}(s,v)$ be the length of the shortest path from s to v that uses exactly i edges, or ∞ if there's no such path.



What is...

$\text{dist}^{(0)}(5,3)?$

$\text{dist}^{(1)}(5,3)?$

$\text{dist}^{(2)}(5,3)?$

$\text{dist}^{(3)}(5,3)?$

$\text{dist}^{(4)}(5,3)?$


Recurrence for SSSP

Key Definition. Let $\text{dist}^{(i)}(s,v)$ be the length of the shortest path from s to v that uses exactly i edges, or ∞ if there's no such path.

End goal: $\min_{i \leq n-1} \text{dist}^{(i)}(s,v)$ for each v .

$\text{dist}^{(i)}(s,v) =$

Base case(s):



The recurrence will be similar in structure to Professor Υ's first attempt

Let's follow the DP Recipe

This is called the **Bellman-Ford Algorithm**
(Bellman, Ford, Moore, Shimbel, 1955)

The DP Recipe

1. Write recurrence ←

usually the trickiest part
2. Size of table: How many dimensions? Range of each dimension?
3. What are the base cases?
4. To fill in a cell, which other cells do I look at? In which order do I fill the table?
5. Which cell(s) contain the final answer?
6. Running time = (size of table) • (time to fill each entry)
7. To reconstruct the solution (instead of just its size) follow arrows from final answer to base case

Pseudocode for Bellman-Ford

Algorithm SSSP(G, s)

table := 2D-array indexed from 0 to $n-1$ in both dimensions

// first dimension represents vertices v_0, \dots, v_{n-1} , where s is v_0 ,
second dimension represents the number i of edges

table(0,0) = 0 // base case for s

for $k = 1$ to $n-1$:

table($k, 0$) = ∞ // rest of base cases

for $i = 1$ to $n-1$: // for each number i of edges

 for $k = 0$ to $n-1$: // for each vertex

table(k, i) = $\min_{(v_j, v_k) \in E} \{ \mathbf{table}(j, i-1) + \ell(v_j, v_k) \}$

Return $\min_i \{ \mathbf{table}(k, i) \}$ for each vertex k

Previously we assumed the graph has no negative-weight cycles. **But what if it does?**

With a slight modification, the Bellman-Ford (BF) Algorithm can detect whether or not there is a negative-weight cycle:

Run BF for n more iterations!

I.e calculate $\text{dist}^{(i)}(s,v)$ for all i up to $2n-1$.

Claim: There is a negative-weight cycle (reachable from s) IFF $\min_{i \leq 2n-1} \text{dist}^{(i)}(s,v) < \min_{i \leq n-1} \text{dist}^{(i)}(s,v)$ for some vertex v .

*Actually it suffices to run BF for just 1 more iteration

Faster Algorithms for SSSP

- Bernstein, Nanongkai, Wulff-Nilsen, 2022: $O(m \cdot \log^8 n)$ ← integer weights
 ↑ ↑
My postdoc Thatchaphol Saranurak's
advisor PhD advisor
- Fineman, 2023: $O(mn^{7/8})$ ← any weights
- If no negative weights, Dijkstra's algorithm: $O(m + n \log n)$

Initial idea for solving APSP: Run SSSP from every vertex

That works, but the algorithm you're about to see is faster for dense graphs: $O(n^3)$

Let's try essentially the same recurrence as BF


$$\text{dist}^{(i)}(s,v) = \min_{(u,v) \in E} \{ \text{dist}^{(i-1)}(s,u) + \ell(u,v) \}$$

$$\text{Base case: } \text{dist}^{(0)}(s,v) = \begin{cases} 0 & \text{if } s = v \\ \infty & \text{otherwise} \end{cases}$$

The only difference:

there are 3 “free” variables: s , v , and $i \Rightarrow 3\text{D DP table!}$

The DP table




Running time:

But there's a cool trick that allows each cell to only look at 3 other cells...

Key idea: Impose an arbitrary ordering on the vertices and consider paths that **only use the first i vertices** in the ordering.

Arbitrary ordering of vertices: v_1, v_2, \dots, v_n .

Key Definition: Let $\text{dist}^{[i]}(s,t)$ be the shortest path from s to t with internal vertices **only** in $\{v_1, \dots, v_i\}$, or ∞ if no such path.



What is...

$\text{dist}^{[0]}(5,3)?$

$\text{dist}^{[1]}(5,3)?$

$\text{dist}^{[2]}(5,3)?$

$\text{dist}^{[3]}(5,3)?$

$\text{dist}^{[4]}(5,3)?$

Recurrence

Arbitrary ordering of vertices: v_1, v_2, \dots, v_n .

Key Definition: Let $\text{dist}^{[i]}(s,t)$ be the shortest path from s to t with internal vertices **only** in $\{v_1, \dots, v_i\}$, or ∞ if no such path.


Base case: $\text{dist}^{[0]}(s,t) = \begin{cases} 0 & \text{if } s = t \\ \ell(s,t) & \text{if } (s,t) \in E \\ \infty & \text{otherwise} \end{cases}$

$\text{dist}^{[i]}(s,t) = \min\{ \quad , \quad \}$

$\underbrace{\hspace{10em}}$
Case 1 (“**Lose it!**”):
shortest st-path
doesn’t contain v_i

$\underbrace{\hspace{10em}}$
Case 2 (“**Use it!**”):
shortest st-path
contains v_i

Let's Follow the DP Recipe



This is called the **Floyd-Warshall Algorithm** (1962)

Pseudocode for Floyd–Warshall

Algorithm APSP(G)

table := 3D-array ($1..n, 1..n, 0..n$)

// first two dimensions represent vertices v_1, \dots, v_n ,

third dimension represents restricting to the first i internal vertices

for $i = 1$ to n :

 for $j = 1$ to n :

$\text{table}(i, j, 0) = \ell(v_i, v_j)$ // base case

for $i = 1$ to n :

 for $j = 1$ to n :

 for $k = 1$ to n :

$\text{table}(i, j, k) = \min\{\text{table}(i, j, k-1), \text{table}(i, k, k-1) + \text{table}(k, j, k-1)\}$

Return **table**(i, j, n) for all i, j


A question to ponder...

Why doesn't the trick of computing $\text{dist}^{[i]}(s,t)$ give a faster-than-BF algorithm for SSSP too?

Progress on APSP since Floyd-Warshall

Author	Runtime	Year
Fredman	$n^3 \log \log^{1/3} n / \log^{1/3} n$	1976
Takaoka	$n^3 \log \log^{1/2} n / \log^{1/2} n$	1992
Dobosiewicz	$n^3 / \log^{1/2} n$	1992
Han	$n^3 \log \log^{5/7} n / \log^{5/7} n$	2004
Takaoka	$n^3 \log \log^2 n / \log n$	2004
Zwick	$n^3 \log \log^{1/2} n / \log n$	2004
Chan	$n^3 / \log n$	2005
Han	$n^3 \log \log^{5/4} n / \log^{5/4} n$	2006
Chan	$n^3 \log \log^3 n / \log^2 n$	2007
Han, Takaoka	$n^3 \log \log n / \log^2 n$	2012
Williams	$n^3 / \exp(\sqrt{\log n})$	2014

This is wild!



Conclusion: Maybe $O(n^{2.99})$ is impossible?

Maybe $O(n^{2.99})$ is impossible?

Either **ALL** of the following have $O(n^{<3})$ time algorithms or **NONE** of them do: (Virginia Vassilevska Williams, Ryan Williams, 2010)

1. APSP
2. Minimum Weight Triangle
3. Metricity
4. Minimum Cycle
5. Distance Product
6. Second Shortest Path
7. Replacement Paths
8. Negative Triangle Listing

...



A curious **open** problem:
“The Not Shortest Path problem”

Given a directed graph with positive edge weights and a pair s, t of vertices, is there a polynomial time algorithm to find a simple path from s to t that is **NOT** shortest?