



More Randomization:

Load Balancing and Fingerprinting



Load Balancing

Imagine Google assigning search queries to servers



Goal:

- No server gets too many clients
- Each client is assigned to a server *without knowledge of the allocation of other clients to servers*

Strategy: Assign each client to a server uniformly at random!

Let's see how well this does...

This is often formulated as:

“Balls and Bins”

n balls (clients)



k bins (servers)

Each ball goes into a random bin.

Question: How many balls in the fullest bin?


For simplicity, we will analyze the case where $n=k$.

We will prove: With prob $\geq 1-1/n$, fullest bin has $O(\log n)$ balls.

This is often formulated as:

“Balls and Bins”

n balls (clients)



First, let's calculate the expected number of balls per bin:

Let X_j be the number of balls in bin j .

Let X_{ij} be an indicator r.v. for whether ball i is in bin j .


Observation: $X_j = \sum_{i=1}^n X_{ij}$

So, $E[X_j] = E\left[\sum_{i=1}^n X_{ij}\right] = \sum_{i=1}^n E[X_{ij}] =$

This is often formulated as:

“Balls and Bins”

n balls (clients)



Our goal is to bound the probability that the fullest bin has “many” balls

Last lecture, to bound probabilities we used *Markov's inequality*.




It turns out that won't suffice here!

We need a stronger bound, that will use the fact that X_j is a sum of independent r.v.'s.

Sum of **Independent** r.v.'s is predictable

Flip a coin n times. What fraction of flips are heads?

(Number of head flips is a sum of independent indicator r.v.'s for whether the i^{th} flip is heads)



Chernoff Bounds

(we won't prove)

How do these
compare to Markov?

Let Y_1, \dots, Y_k be independent r.v.'s taking values in the range $[0, 1]$.

Let $Y = \sum_{i=1}^k Y_i$. Let $\mu = E[Y]$.

woohoo! I'm
in the bounds



“Large Deviation” bound:

For any $\lambda \geq 1$:

$$\Pr[Y - \mu \geq \lambda\mu] \leq e^{-\lambda\mu/3}$$

“Small Deviation” bounds:

For any $\lambda \in [0, 1]$:

$$\Pr[Y - \mu \geq \lambda\mu] \leq e^{-\lambda^2\mu/3}$$

$$\Pr[Y - \mu \leq -\lambda\mu] \leq e^{-\lambda^2\mu/3}$$


this one we won't
use in class, but
you may use it on
the HW

Another useful tool:

Union Bound

- for arbitrary events A, B

$$\Pr[A \cup B] \leq \Pr[A] + \Pr[B]$$





- More generally, for arbitrary events E_1, \dots, E_n

$$\Pr[\cup_i E_i] \leq \sum_i \Pr[E_i]$$

Now back to Balls and Bins

n balls (clients)

Let's apply the new tools in our pocket (Chernoff and union bounds)



Recall $\mathbf{X_j}$ is the number of balls in bin j . Earlier we showed: $E[\mathbf{X_j}] = 1$

Recall $\mathbf{X_{ij}}$ is an indicator r.v. for whether ball i is in bin j .



Observation: $\mathbf{X_j} = \sum_{i=1}^n \mathbf{X_{ij}}$.

$\Rightarrow \mathbf{X_j}$ is the sum of independent r.v.'s with values in $[0,1]$ so we can apply Chernoff.

Now back to Balls and Bins

n balls (clients)

Markov instead of Chernoff gives prob $1/\log n$, which becomes $n/\log n$ after union bound. Not useful!



Apply Chernoff: $\Pr[X_j \geq 1 + 6 \ln n] \leq$

Apply Union: $\Pr[\text{Exists } j \text{ such that } X_j \geq 1 + 6 \ln n] \leq$

\Rightarrow Original goal: with prob $\geq 1 - 1/n$, fullest bin has $O(\log n)$ balls.


Fingerprinting

The scenario:

- You download a large file from a *untrusted* remote server
- The original file is from your friend
- You want to check that the version you downloaded hasn't be tampered with.
- The file is large so your friend can't send the whole thing to you directly, but they can send you a small “**fingerprint**” to help verify the authenticity



λ -Productions is proud to present a visiting cast of characters...




Beforehand, Alice and Bob agree on a **protocol** for how Alice will choose **M**, given **x**.



Alice has **x** (original file)


Message **M** (the “fingerprint”)



Bob has **y**, **M**





malicious adversary
who knows protocol



y
(downloaded
file)

heh heh I will pick **y** so that
 $\mathbf{x} \pmod{10} \equiv \mathbf{y} \pmod{10}$
but $\mathbf{y} \neq \mathbf{x}$ and Bob will never
know I changed the file



malicious adversary
who knows protocol

y
(downloaded
file)

I will interpret the
entire file **x** as a
number, and send
 $\mathbf{M} = \mathbf{x} \pmod{10}$

I will say " $\mathbf{x} = \mathbf{y}$ " iff
 $\mathbf{x} \pmod{10} \equiv \mathbf{y} \pmod{10}$

Message **M** (the "fingerprint")



Alice has **x** (original file)



Bob has **y**, **M**

Goal: We want message **M** as short as possible, while still
ensuring that no matter how the adversary changes the file,
Bob can check if $\mathbf{x} = \mathbf{y}$, given **M**.

Deterministic Protocols Don't Work

If $|M| < |x|$, then by the pigeonhole principle, there are two files x_1, x_2 that deterministically cause Alice to send the same message M_{bad} .

If Bob receives the message M_{bad} and downloads the file $y = x_1$, Bob doesn't know if the original file was x_1 or x_2 ! He deterministically says either “ $x = y$ ” (wrong if $x = x_2$), or “ $x \neq y$ ” (wrong if $x = x_1$).

Therefore, all deterministic protocols require $|M| = |x|$ in the worst case.

This is useless because the file is too large to send!

But there is a good **randomized** protocol!



We will show:

There is a randomized protocol such that for every x, y :

- $|M| = O(\log n)$ $\leftarrow n$ is #bits of x
- If $x = y$, then Bob always says “ $x = y$ ”
- If $x \neq y$, then Bob detects that “ $x \neq y$ ” with prob $\geq 90\%$

1st attempt

heh heh I will pick **y** so that



malicious adversary
who knows protocol

y
(downloaded
file)

I will pick a random
number p in $[1..10]$
and send
 $M = (p, x \pmod{p})$

I will say " $x = y$ " iff
 $x \pmod{p} \equiv y \pmod{p}$

Message **M** (the "fingerprint")



Alice has **x** (original file)




Bob has **y**, **M**

Goal: We want message **M** as short as possible, while still ensuring that no matter how the adversary changes the file, Bob can check if $x = y$, given **M**.

2nd attempt

I want to pick y so that $|x - y|$ has as many factors in $[1..n]$ as possible.



malicious adversary
who knows protocol

y
(downloaded
file)

I will pick a random
number p in $[1..n]$
and send
 $M = (p, x \pmod{p})$

I will say " $x = y$ " iff
 $x \pmod{p} \equiv y \pmod{p}$

Message M (the "fingerprint")



Alice has x (original file)



Bob has y, M

Goal: We want message M as short as possible, while still ensuring that no matter how the adversary changes the file, Bob can check if $x = y$, given M .

Adversary wants $|x-y|$ to have many factors, we want few factors.

How many factors does a number have?

Exponential in the number of prime factors.

That seems like a lot...

Insight: Alice picks only from the set of prime numbers!

Let's see why this works...

The actual protocol



I want to pick **y** so that $|x - y|$ has as many **prime** factors as possible.
But no matter what **y** I pick I can't seem to fool Bob...

I will pick p randomly from the first $10n$ **prime** numbers and send $M = (p, x \pmod{p})$



Alice has **x** (original file)

malicious adversary
who knows protocol



y
(downloaded file)

I will say " $x = y$ " iff $x \pmod{p} \equiv y \pmod{p}$

Message **M** (the "fingerprint")



Bob has **y**, **M**

Goal: We want message **M** as short as possible, while still ensuring that no matter how the adversary changes the file, Bob can check if $x = y$, given **M**.

Our goal is to show:

For all x, y the protocol is such that:

- $|M| = O(\log n)$ $\leftarrow n$ is #bits of x
- If $x = y$, then Bob always says “ $x = y$ ”
- If $x \neq y$, then Bob detects that “ $x \neq y$ ” with prob $\geq 90\%$

$M = (p, x \pmod p)$ where p is among the first $10n$ primes.

Question: How big is the k^{th} prime number?

Answer: $O(k \log k)$ (complicated proof from number theory)

So, $p = O(n \log n)$.

$\Rightarrow |M| = O(\text{\#bits in } p) = O(\log(n \log n)) = O(\log n)$.

Our goal is to show:

For all x, y the protocol is such that:

- $|M| = O(\log n)$ $\leftarrow n$ is #bits of x
- If $x = y$, then Bob always says “ $x = y$ ”
- If $x \neq y$, then Bob detects that “ $x \neq y$ ” with prob $\geq 90\%$

If $x \neq y$, then Bob wrongly answers “ $x = y$ ” if $x \pmod{p} \equiv y \pmod{p}$,
i.e. if p divides $|x - y|$.

Question: How many primes divide $|x - y|$?

Answer: $\leq n$. Why?

Alice chooses from $10n$ primes, and $\leq n$ of them cause Bob to wrongly answer “ $x = y$ ”. Thus, Bob is right with prob $\geq 90\%$.

What if we wanted Bob to succeed 99% of the time
instead of 90%?

Some takeaways from today

- Union & Chernoff bounds: used frequently in probabilistic analysis
 - Chernoff: “sum of independent r.v.’s is predictable”
 - Union: for calculating the prob that no “bad” event happens
- Randomization is necessary for frequently arising situations:
 - Load Balancing: allocating clients to servers
 - Fingerprinting: verifying authenticity of file from remote server