

$$k \geq 1, b > 1, d \geq 0, w \geq 0$$

$$T(n) = k \cdot T(n/b) + O(n^d)$$

$$= \begin{cases} O(n^d) & \text{if } k < b^d \\ O(n^d \log n) & \text{if } k = b^d \\ O(n^{\log_b k}) & \text{if } k > b^d \end{cases}$$

$$T(n) = kT(n/b) + O(n^d \log^w n)$$

$$\begin{cases} O(n^d \log^w n) & \text{if } k < b^d \\ O(n^d \log^{w+1} n) & \text{if } k = b^d \\ O(n^{\log_b k}) & \text{if } k > b^d \end{cases}$$

D&C: Int multi

- Let's stare at this identity again:

$$xy = (a \cdot 10^{n/2} + b)(c \cdot 10^{n/2} + d) \\ = ac \cdot 10^n + (ad + bc) \cdot 10^{n/2} + bd$$

- Think:**

- Could we write $ad + bc$ in terms of $ac(t_1)$, $bd(t_3)$, and something else that only uses one multiplication (not two)?

$$ad + bc = (a + b)(c + d) - ac - bd$$

D&C: Closest Pair(nlogn)

Sorted by x Sorted by y

ClosestPair($P_1, \dots, P_n, P'_1, \dots, P'_n$): // $n \geq 3$ pts in the plane

if $n \leq 3$ then return min dist among P_1, P_2, P_3 // base case

$(L, R) \leftarrow$ partition points by $P_{n/2}$ // split by median x-coordinate

$\delta_1 \leftarrow$ **ClosestPair**(L (sorted by x), L (sorted by y)) // min dist on left

$\delta_2 \leftarrow$ **ClosestPair**(R (sorted by x), R (sorted by y)) // min dist on right

$\delta_3 \leftarrow$ min distance in δ -strip // details in notes

return $\min\{\delta_1, \delta_2, \delta_3\}$

DP: Runtime = Size of table *
time to fill every entry

DP: Task selection(nlogn if sorted time)

OPT has J_n (use it!) \leftrightarrow $OTV(J_1, \dots, J_n) = \text{val}(J_n) + OTV(J_1, \dots, J_i)$

J_i is the last interval that doesn't overlap with J_n

OPT doesn't have J_n (lose it!) \leftrightarrow $OTV(J_1, \dots, J_n) = OTV(J_1, \dots, J_{n-1})$

"Optimal Task Value" i.e. the value of the optimal solution

DP: Longest Increasing subseq(n^2)

Define $CLIS(S[1..N])$ to be the longest of the increasing subsequences of $S[1..N]$ that contain $S[N]$.

Define $L(i)$ to be the length of $CLIS(S[1..i])$

Then, we have:

$$L(i) = 1 + \max\{L(j) : 0 < j < i \text{ and } S[j] < S[i]\}$$

DP: LCS(MN)

• **Def:** $LCS(i, j)$ = length of a LCS of $X[1..i]$ and $Y[1..j]$.

• $i = 0$ means X is the empty string

• $j = 0$ means Y is the empty string

• **Goal:** return $LCS(m, n)$

$$\text{We have: } LCS(i, j) = \begin{cases} 0 & \text{if } i = 0 \text{ or } j = 0 \\ 1 + LCS(i-1, j-1) & \text{if } X[i] = Y[j] \\ \max\{LCS(i-1, j), LCS(i, j-1)\} & \text{if } X[i] \neq Y[j] \end{cases}$$

DP:0-1(nW)

- The recurrence:**

$$Knapsack(\{t_1, \dots, t_n\}, W) = \max \left\{ \begin{array}{l} Knapsack(\{t_1, \dots, t_{n-1}\}, W) \\ v_n + Knapsack(\{t_1, \dots, t_{n-1}\}, W - w_n) \end{array} \right.$$

- The base cases:**

– $Knapsack(\emptyset, W') = 0$ for all W'

– $Knapsack(\{t_1, \dots, t_i\}, 0) = 0$ for all i

Valid only when $W \geq w_n$

DP:Longest Pallindromic substrng

$$PAL(i, j) = \begin{cases} X[i] = X[j] & j \leq i + 1 \\ X[i] = X[j] \text{ and } PAL(i+1, j-1) & j > i + 1 \end{cases}$$

GraphDP: Singlesource Shortest path(BF)

Definition

- $dist^{(i)}(s, t)$ = " i -hop distance from s to t "
shortest length of an $s \rightarrow t$ path using exactly i edges, or ∞ if there's no such path
- $dist^{(\leq i)}(s, t)$ = "at-most- i -hop distance from s to t "
shortest length of an $s \rightarrow t$ path using at most i edges

Lemma:

In n -node graph without neg-length cycles,

$$dist^{(\leq n-1)}(s, t) = dist(s, t)$$

$$- dist^{(i)}(s, v) \leftarrow \min_u dist^{(i-1)}(s, u) + \ell(u, v)$$

GraphDP: All pair BF

$$dist^{(\leq i)}(s, t) = \min_x dist^{(\leq i/2)}(s, x) + dist^{(\leq i/2)}(x, t)$$

GraphDP: Floyd

Definition

$dist^{[i]}(s, t)$ is the "middle-restricted distance:"
Shortest length of an $s \rightarrow t$ path that only uses $\{v_1, \dots, v_i\}$ as intermediate vertices (but s, t can be anything)

$$dist^{[k]}(s, t) = \min \left\{ \begin{array}{l} dist^{[k-1]}(s, t) \\ dist^{[k-1]}(s, v_k) + dist^{[k-1]}(v_k, t) \end{array} \right.$$

- Bellman-Ford (naïve method):**

$$- O(mn^2) \text{ time}$$

$$|E||V|^2$$

- Bellman-Ford (with path-doubling):**

$$- O(n^3 \log n) \text{ time}$$

$$|V|^3 / \log |V|$$

- Floyd-Warshall (next):**

$$- O(n^3) \text{ time}$$

$$|V|^3$$

Definition 34 (Tree #1) An undirected graph G is a *tree* if it is connected and acyclic (i.e., has no cycle).

A graph is *connected* if for any two vertices there is a path between them. A *cycle* is a nonempty sequence of adjacent edges that starts and ends at the same vertex.

Definition 35 (Tree #2) An undirected graph G is a tree if it is *minimally connected*, i.e., if it is connected, and removing any single edge causes it to become disconnected.

Definition 36 (Tree #3) An undirected graph G is a tree if it is *maximally acyclic*, i.e., if it has no cycle, and adding any single edge causes it to have a cycle.

A decision problem can be thought of as

$$f: \Sigma^* \rightarrow \{\text{No}, \text{Yes}\}$$

or equivalently as a language

$$L \subseteq \Sigma^*$$

$$L = \{x \in \Sigma^* : f(x) = \text{Yes}\} \quad f(x) = \begin{cases} \text{Yes if } x \in L \\ \text{No if } x \notin L \end{cases}$$

Let M be a DFA, using alphabet Σ .

M accepts some strings in Σ^* and rejects the rest.

Definition: $L(M) = \{x \in \Sigma^* : M \text{ accepts } x\}$

Called the "language decided by M ".

If L is a language,

we say M decides L if $L(M) = L$.

A deterministic finite automaton is a 5-tuple:

$$M = (Q, \Sigma, \delta, q_0, F)$$

① Q is a nonempty finite set of states,

② Σ is an alphabet,

③ $\delta: Q \times \Sigma \rightarrow Q$ is the state-transition function,

④ $q_0 \in Q$ is the start state,

⑤ $F \subseteq Q$ is the set of accepting states.

即: graph 的 abstraction. 有几个 states, 哪个 start, 哪些 acc, 使用什么 alphabet, states 间怎么 transit.

Regular Expression

A regular expression = finite expression using the string notations

- Start from finite alphabet.
 - Compose them using concatenation, alternation " $|$ " or Kleen star " $*$ ".
- Examples:
- $L(\epsilon) = \{\epsilon\}$
 - $L(ab^*|ba^*) = \{a, ab, abb, \dots\} \cup \{b, ba, baa, \dots\}$

Theorem (RegExp = DFA):

L is defined by regular expression $\Leftrightarrow L$ is decided by a DFA

These languages are called regular languages.

A Turing machine is a 7-tuple:

$$M = (Q, \Gamma, \Sigma, \delta, q_{\text{start}}, q_{\text{accept}}, q_{\text{reject}})$$

- Q = set of states
- Σ = the input alphabet (typically $\{0,1\}$ but not always)
- \perp = the blank symbol
- Γ = the tape alphabet where generally $\Gamma = \Sigma \cup \{\perp\}$
- $q_{\text{start}} \in Q$, = the initial state
- $F = \{q_{\text{accept}}, q_{\text{reject}}\} \subseteq Q$, = the set of final states (one accepting state and one rejecting state)
- $\delta: (Q \setminus F) \times \Gamma \rightarrow Q \times \Gamma \times \{L, R\}$ = the transition function

All of these sets are finite



(现 state, symbol) \rightarrow (下 state, symbol, pointer/L/R)
(read) (write)

Definition: A Turing Machine M recognizes a language L if it:

- accepts every string in L , and
- rejects OR loops for every string not in L .

Definition: A Turing Machine M decides a language L if it:

- accepts every string in L , and
- rejects every string not in L (and never loops forever)

A language L is decidable if there is a TM that decides L .

Fact: If a language L is decidable, then the complement of L is also decidable. Why?

Fact: If languages L_1 and L_2 are both decidable, then $L_1 \cap L_2$ and $L_1 \cup L_2$ are both decidable. Why?

Theorem: Turing machines can simulate each of these models (after formalization) and all known models of computation.

Lemma 74 The (infinite) set $\{0, 1\}^*$ of binary strings is countable.

Lemma 76 The (infinite) set \mathcal{M} of Turing machines is countable.

Proof 77 The key observation is that any Turing machine M has a finite description, and hence can be encoded as a (finite) binary string $\langle M \rangle \in \Sigma^*$ in some unambiguous way. To see this, notice that all the components of the seven-tuple are finite: the alphabets Σ, Γ , the set of states Q and the special states $q_{\text{start}}, q_{\text{acc}}, q_{\text{rej}}$, and the transition function δ . In particular, δ has a finite domain and codomain, so we can encode its list of input/output pairs as a (finite) binary string.

There is an undecidable language $A \subseteq \Sigma^* = \{0, 1\}^*$.

	ϵ	0	1	00	01	10	11	000	...
M_0	yes	no	yes	yes	no	no	no	no	
M_1	yes	yes	no	no	no	no	no	yes	
M_2	yes	no	no	no	yes	no	no	yes	
M_3	yes	no	yes	no	yes	yes	no	no	
M_4	yes	no	yes	no	yes	no	no	yes	
M_5	yes	yes	no	no	no	no	yes	yes	
M_6	no	yes	no	no	yes	no	no	no	
M_7	yes	no	no	yes	no	yes	no	yes	

Specifically, an interpreter U takes two inputs: (1) source code $\langle M \rangle$, and (2) string x .

U simulates the execution of M on input x :

- M accepts $x \Rightarrow U$ accepts $\langle M \rangle, x$
- M rejects $x \Rightarrow U$ rejects $\langle M \rangle, x$
- M loops on $x \Rightarrow U$ loops on $\langle M \rangle, x$

This is called the Universal Turing Machine (and it does exist)

$L_{\text{DAPED}} = \{\langle M \rangle : M \text{ is a TM and } M(\langle M \rangle) \text{ does not accept}\}$

$L_{\text{ACC}} = \{\langle M \rangle, x : M \text{ is a Turing machine and } M(x) \text{ accepts}\}$

$L_{\text{HALT}} = \{\langle M \rangle, x : M \text{ halts on input } x\}$

$M_{\text{ACC}}(\langle M \rangle, x)$:

Run M_{HALT} on $\langle M \rangle, x$
If M_{HALT} rejects, reject
If M_{HALT} accepts, run M on x
If M accepts, accept
If M rejects, reject

$$L_{\text{ACC}} \leq_T L_{\text{HALT}}$$

$M_{\text{BARBER}}(\langle M \rangle)$:

Run M_{ACC} on $\langle M \rangle, \langle M \rangle$.
If M_{ACC} accepts, then reject.
If M_{ACC} rejects, then accept.

$$L_{\text{BARBER}} \leq_T L_{\text{ACC}}$$

$M_{\text{HALT}}(\langle M \rangle, x)$:

Let M_x be a TM that ignores its input and runs $M(x)$
Run $M_{\text{E-HALT}}(\langle M_x \rangle)$ and answer as $M_{\text{E-HALT}}$

Run $M(x)$ and answer as M