


Greedy Algorithms




A question to ponder...

Why doesn't the trick of computing $\text{dist}^{[i]}(s,t)$ give a faster-than-BF algorithm for SSSP too?

Progress on APSP since Floyd-Warshall

Author	Runtime	Year
Fredman	$n^3 \log \log^{1/3} n / \log^{1/3} n$	1976
Takaoka	$n^3 \log \log^{1/2} n / \log^{1/2} n$	1992
Dobosiewicz	$n^3 / \log^{1/2} n$	1992
Han	$n^3 \log \log^{5/7} n / \log^{5/7} n$	2004
Takaoka	$n^3 \log \log^2 n / \log n$	2004
Zwick	$n^3 \log \log^{1/2} n / \log n$	2004
Chan	$n^3 / \log n$	2005
Han	$n^3 \log \log^{5/4} n / \log^{5/4} n$	2006
Chan	$n^3 \log \log^3 n / \log^2 n$	2007
Han, Takaoka	$n^3 \log \log n / \log^2 n$	2012
Williams	$n^3 / \exp(\sqrt{\log n})$	2014

This is wild!



Conclusion: Maybe $O(n^{2.99})$ is impossible?

Maybe $O(n^{2.99})$ is impossible?

Either **ALL** of the following have $O(n^{<3})$ time algorithms or **NONE** of them do: (Virginia Vassilevska Williams, Ryan Williams, 2010)

1. APSP
2. Minimum Weight Triangle
3. Metricity
4. Minimum Cycle
5. Distance Product
6. Second Shortest Path
7. Replacement Paths
8. Negative Triangle Listing
- ...



Greedy Algorithms

Pick the best choice **NOW**.

Prove you end up with an optimal solution.



Proof Technique: Induction + “Exchange” argument


Warning: Greed is generally bad!



Greed

Divide and conquer

Dynamic programming



- Fast
- doesn't work for most problems

- Often slower than greed and faster than DP
- works when solutions to disjoint subproblems can be combined into final solution

- Generally slower (but still usually efficient)
- applies to many problems

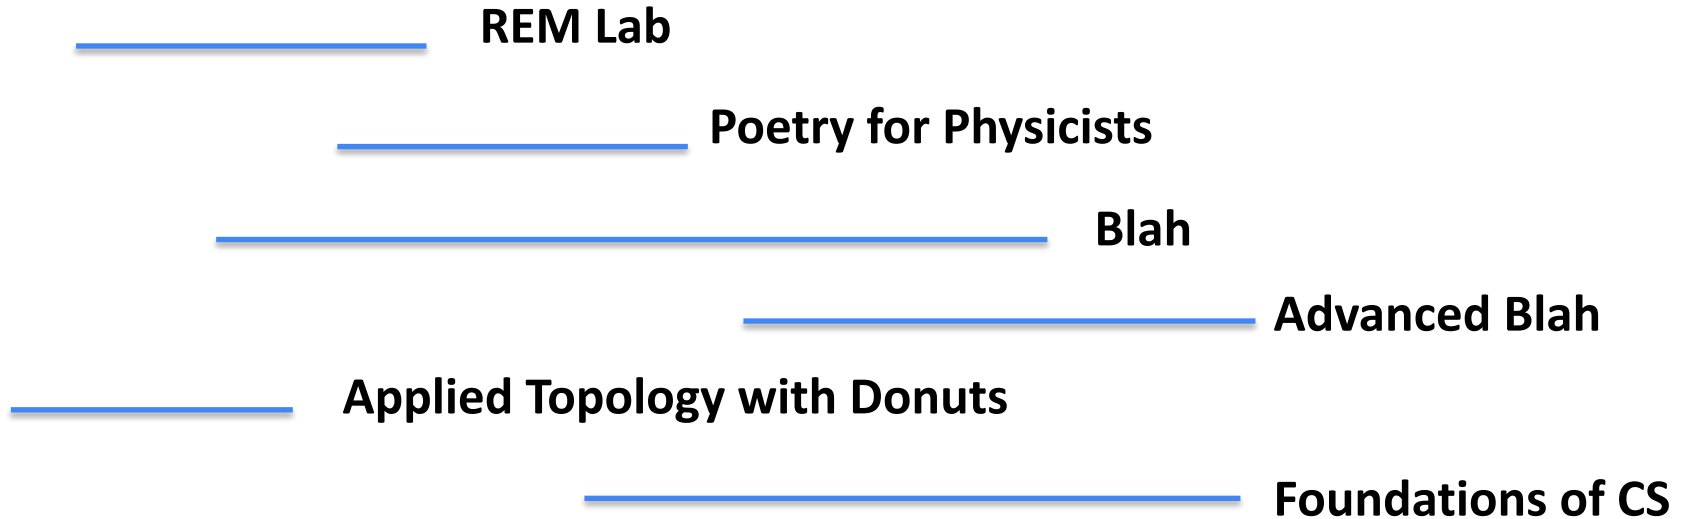
⋮

harder problems where none of these methods apply (coming soon)

But sometimes greed can be good...

Unweighted Course Registration



(aka Task Scheduling)




Goal: Choose a **largest*** possible set of non-intersecting courses (there may be many optimal solutions, we just seek one!)

*We do not endorse this course registration strategy

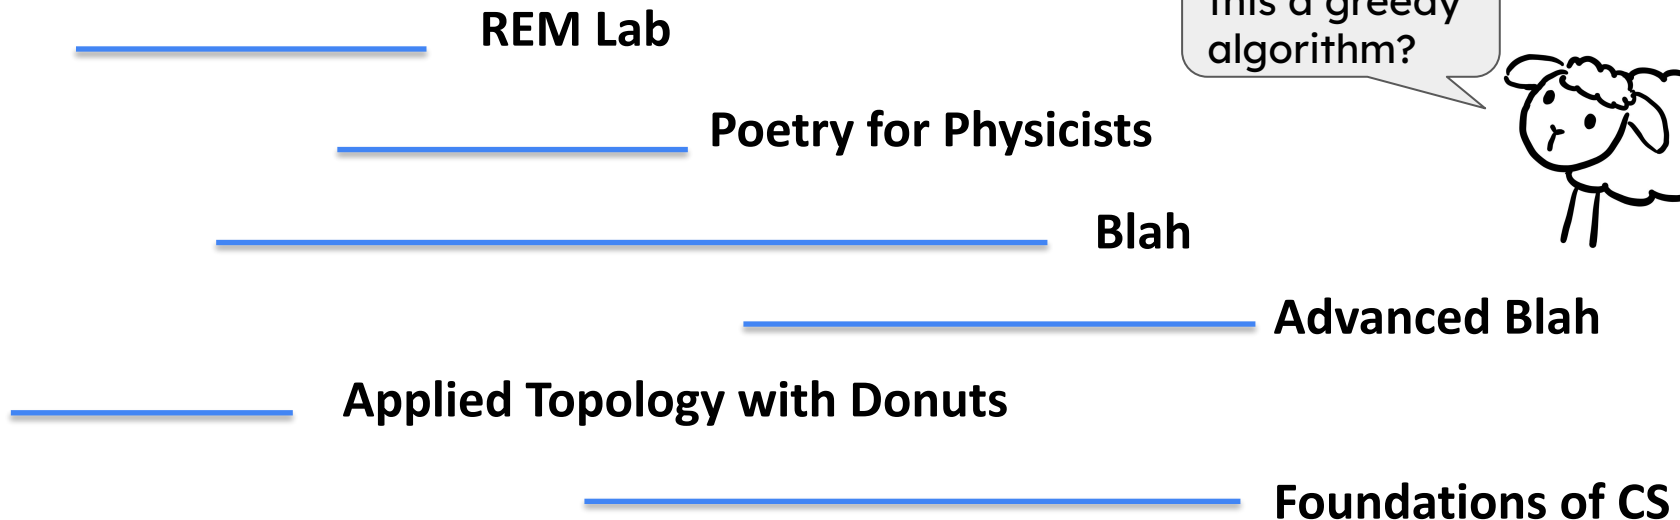
I have not 1, not 2, but 3
greedy algorithms!




Professor Υ 's Greedy Algorithms



Attempt 1: Choose the **shortest interval** (breaking ties arbitrarily), take it, remove overlaps, recurse on remaining problem.




What makes
this a greedy
algorithm?



Counterexample:


Professor Υ 's Greedy Algorithms



Attempt 2: Choose the **interval that starts earliest** (breaking ties arbitrarily), take it, remove overlaps, recurse on remaining problem.

Counterexample:

Professor Υ 's Greedy Algorithms




Attempt 3: Choose the **interval that overlaps with the fewest other intervals** (breaking ties arbitrarily), take it, remove overlaps, recurse on remaining problem.

Counterexample:

A Correct Greedy Algorithm: “Earliest Ending Time (EET)”

- Sort the intervals by *ending time*
- Greedily take the **interval I that ends first** (break ties arbitrarily)
- Remove the intervals that overlap with the one just selected
- Recurse on remaining problem



A Correct Greedy Algorithm: “Earliest Ending Time (EET)”

- Sort the intervals by *ending time*
- Greedily take the **interval I that ends first** (break ties arbitrarily)
- Remove the intervals that overlap with the one just selected
- Recurse on remaining problem




A Correct Greedy Algorithm: “Earliest Ending Time (EET)”

- Sort the intervals by *ending time*
- Greedily take the **interval I that ends first** (break ties arbitrarily)
- Remove the intervals that overlap with the one just selected
- Recurse on remaining problem



Let's see the big
idea of the proof of
correctness first
and then the formal
proof afterwards





A Correct Greedy Algorithm: “Earliest Ending Time (EET)”

Key Claim: The interval I that ends first is a **safe** choice i.e. it is in some optimal solution.

Why? Consider an optimal solution OPT . Let I_{OPT} be the interval that ends first in OPT .

- I_{OPT} ends at least as late as I .
- All other intervals in OPT start after I_{OPT} ends, and thus after I ends.
- Thus, we can take OPT and **exchange** I_{OPT} for I and get a valid solution of the same size!

I



Formal Proof of Correctness by Induction

intervals in order of addition (so, increasing end times)

Let I_1, I_2, I_3, \dots be the output of the EET algorithm


Goal: Prove that for all k , $I_1, I_2, I_3, \dots, I_k$ is in some optimal solution.

Proof by induction on k :

Base case: $k=0$.

Inductive hypothesis: Suppose $I_1, I_2, I_3, \dots, I_k$ is in some optimal solution $I_1, I_2, I_3, \dots, I_k, I'_{k+1}, I'_{k+2}, \dots$

Inductive step: Goal: show $I_1, I_2, I_3, \dots, I_{k+1}$ is in some optimal solution.



Formal Proof of Correctness by Induction


Inductive step: Goal: show $I_1, I_2, I_3, \dots, I_{k+1}$ is in some optimal solution.

By the inductive hypothesis, there exists an optimal solution:

$$\text{OPT} = I_1, I_2, I_3, \dots, I_k, I'_{k+1}, I'_{k+2}, \dots$$

Use an **exchange** argument as before!

- I'_{k+1} ends at least as late as I_{k+1} .
- All other intervals I'_{k+2}, \dots start after I'_{k+1} ends, and thus after I_{k+1} ends.
- Thus, we can take OPT and **exchange** I'_{k+1} for I_{k+1} and get a valid solution of the same size!



General Strategy commonly used for analyzing greedy algorithms:

Proof by induction using an “**exchange**” argument

The idea: Show that we can transform any **optimal solution** into the **solution given by our algorithm** by **exchanging** each piece of it out one-by-one without increasing the cost.

Key part of proof: Show that my greedy choice is **safe** i.e. it is in some optimal solution.


Induction just formalizes the idea that *each successive choice* is **safe**.

Another problem where greed is good...


A Highway Problem

Input: an undirected graph with positive edge weights
e.g. a set of cities and distances between them

Output: minimum total length of highway to connect all cities
i.e. it should be possible to drive from any city to any other
using just the highways



a and **d** are
Ann Arbor and
Detroit



Another problem where greed is good...


A Highway Problem

Claim: Solution is a tree.

↙ This fact will be useful later too

Why? If a connected graph has a cycle, we can delete an arbitrary edge from the cycle and still have a connected graph.

Definition (review): A tree is a connected graph with no cycles.




The Highway Problem is commonly known as:

Minimum Spanning Tree (MST)

Given a graph G , a **spanning tree** is a subgraph of G that is spanning (uses all vertices) and is a tree.

MST problem: Given an undirected graph with positive edge weights, find a minimum weight spanning tree.




Minimum Spanning Tree (MST)


Template for greedy algorithm:

Greedly pick an edge and add it to our MST. Repeat.

But which edge do we pick?



We need to pick a **safe** edge: an edge that appears in *some* optimal solution with all of the previously picked edges.



Kruskal's Algorithm (1956)

Pick the minimum weight edge!



Kruskal(G): // G is a weighted, undirected graph

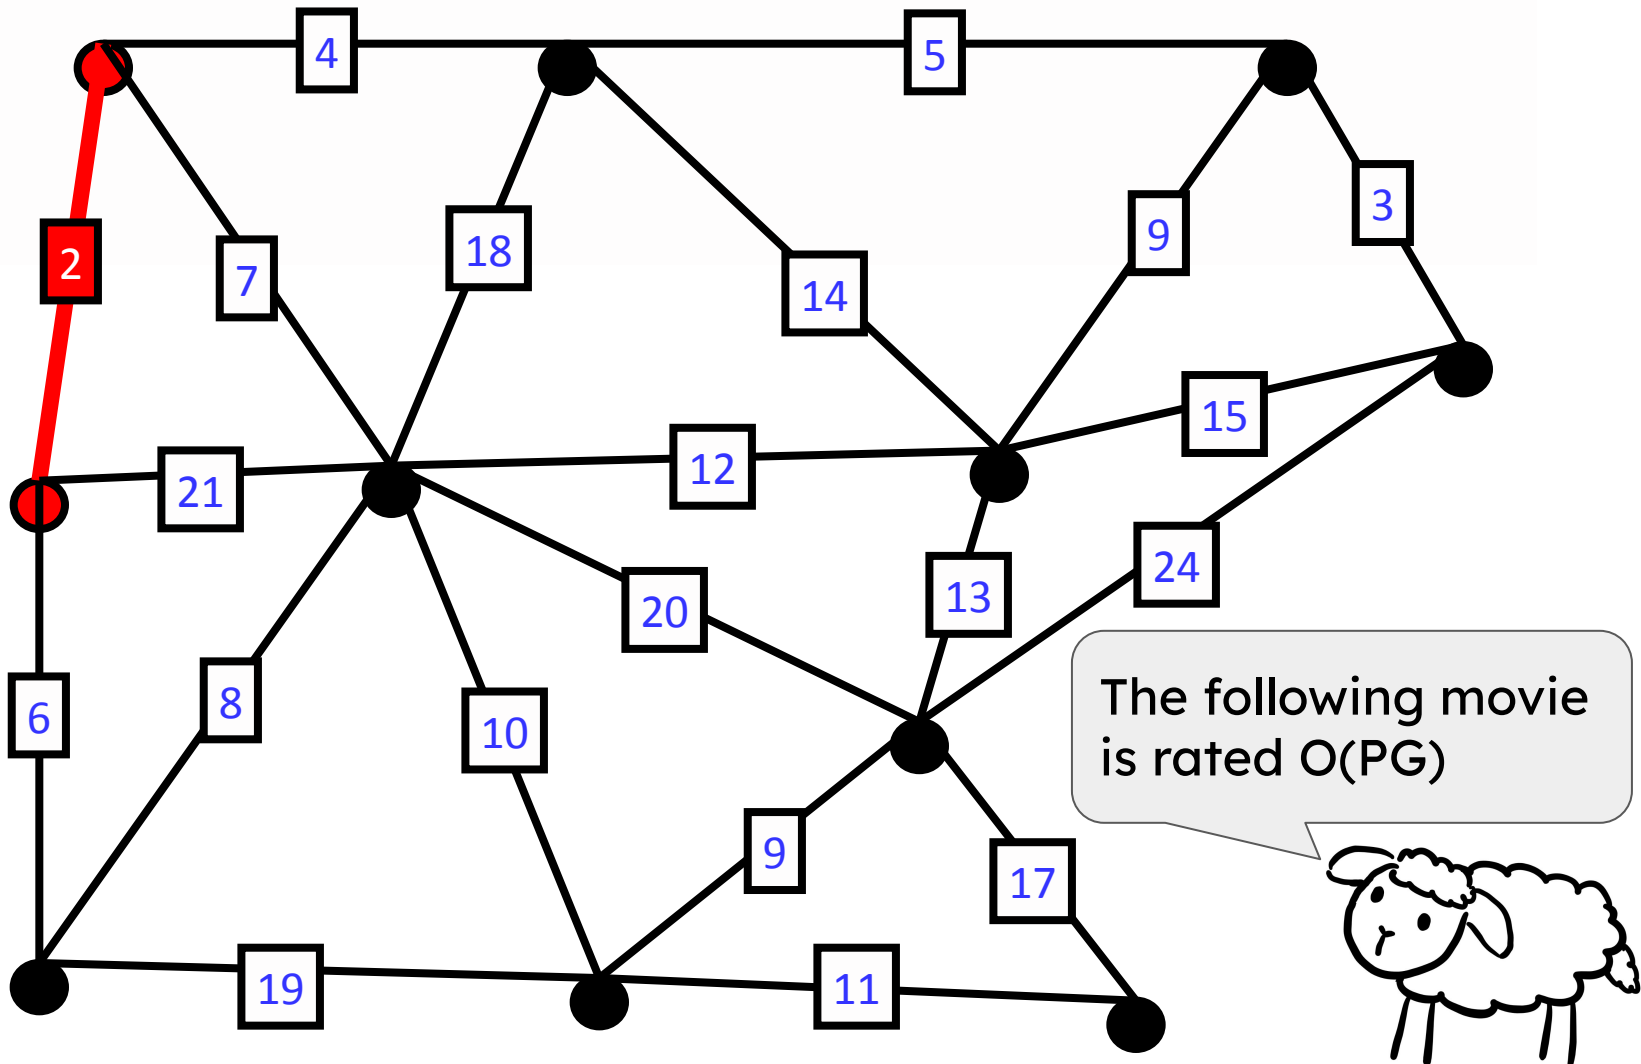
$T \leftarrow \emptyset$ // invariant: T has no cycles

for each edge e in *increasing order of weight*:


if $T + e$ is acyclic: $T \leftarrow T + e$

return T


Sorted weights: **2**,3,4,5,6,7,8,9,9,10,11,12,13,14,15,17,18,19,20,21,24




Sorted weights: 2, **3**, 4, 5, 6, 7, 8, 9, 9, 10, 11, 12, 13, 14, 15, 17, 18, 19, 20, 21, 24




Sorted weights: 2,3,4,5,6,7,8,9,9,10,11,12,13,14,15,17,18,19,20,21,24




Sorted weights: 2,3,4,5,6,7,8,9,9,10,11,12,13,14,15,17,18,19,20,21,24




Sorted weights: 2,3,4,5,**6**,7,8,9,9,10,11,12,13,14,15,17,18,19,20,21,24




Sorted weights: 2,3,4,5,6,**7**,8,9,9,10,11,12,13,14,15,17,18,19,20,21,24




Sorted weights: 2,3,4,5,6,7,**8**,9,9,10,11,12,13,14,15,17,18,19,20,21,24




Sorted weights: 2,3,4,5,6,7,8,9,9,10,11,12,13,14,15,17,18,19,20,21,24




Sorted weights: 2,3,4,5,6,7,8,9,**9**,10,11,12,13,14,15,17,18,19,20,21,24




Sorted weights: 2,3,4,5,6,7,8,9,9,10,11,12,13,14,15,17,18,19,20,21,24




Sorted weights: 2,3,4,5,6,7,8,9,9,10,**11**,12,13,14,15,17,18,19,20,21,24



Sorted weights: 2,3,4,5,6,7,8,9,9,10,11,**12**,13,14,15,17,18,19,20,21,24





Kruskal's Algorithm: Correctness

Why does Kruskal's algorithm return a **spanning tree**?

1. Why is it a **tree**?
2. Why is it **spanning** (i.e. connects all vertices)?

Now we need to show it returns a **minimum** spanning tree.

Again, we'll use induction with an **exchange** argument.

Setting up the Induction

edges in order of addition (so, **non-decreasing** weights)

Let $T = e_1, e_2, e_3, \dots$ be the output of Kruskal's algorithm

Goal: Prove that for all k , the edge set $e_1, e_2, e_3, \dots, e_k$ is in some MST.

Proof by induction on k :

Base case: $k=0$.

Inductive hypothesis: Suppose the edge set $e_1, e_2, e_3, \dots, e_k$ is in some MST $T' = e_1, e_2, e_3, \dots, e_k, f_1, f_2, \dots$ $\leftarrow f$ edges listed in no particular order

Inductive step: Goal: Show the edge set $e_1, e_2, e_3, \dots, e_{k+1}$ is in some MST.

Inductive step

Inductive step: Goal: Show the edge set $e_1, e_2, e_3, \dots, e_{k+1}$ is in some MST.

By the inductive hypothesis, there exists an MST:

$$T' = e_1, e_2, e_3, \dots, e_k, f_1, f_2, \dots$$

Case 1: $e_{k+1} \in T'$

We are done.

Case 2: $e_{k+1} \notin T'$

Claim: We can perform an edge **exchange**: We can take T' , add e_{k+1} , and remove an edge in f_1, f_2, \dots to get a spanning tree that still has minimum weight.

Proof of Claim:

Which edge f_i do we exchange with e_{k+1} ?

After the exchange, why is the graph still connected?

After the exchange, why is the graph still a tree?

Proof of Claim (continued):

Why did the exchange not increase the weight of the spanning tree?

Want to show: $\text{weight}(\mathbf{e}_{k+1}) \leq \text{weight}(\mathbf{f}_i)$.

Other independent inventors of MST algorithms:

Jarník, Prim, Borůvka



First used for designing electrical networks in 1926!

Today, a number of algorithms for clustering data (useful in machine learning) are variations of Kruskal's algorithm

Running time (we won't prove) of Kruskal's algorithm: $O(m \log n)$
 \Rightarrow need to analyze a data structure for detecting cycles

Best-known running time (Chazelle 2000): $O(m \cdot \alpha(n))$

Open problem: $O(m)$?


Inverse Ackermann
function: an extremely
slow growing function:
 $\alpha(n) \leq 4$ even when n is
particles in known
universe

Seth Pettie and Vijaya Ramachandran (2002) gave an asymptotically
optimal algorithm. But nobody knows how fast it is!



You've just finished a crash course on algorithms!

If you found the material interesting, try:



Next up: **Computability**

- What is a computer?
- What problems can a computer solve?
- What problems can a computer **not** solve?