

More Randomized Algorithms

```
int getRandomNumber()  
{  
    return 4; // chosen by fair dice roll.  
              // guaranteed to be random.  
}
```

When is Randomization Necessary?

my “bonus” lecture



next week



- In the areas of online algorithms / cryptography / games / etc. when the input is hidden, we can often **prove** randomization is necessary.
- In the “standard” setting (input given upfront), whether or not randomization is necessary is a **big open problem**. Most researchers seem to believe every randomized algorithm can be derandomized.
- When randomization isn’t necessary, it’s still useful for getting simpler and/or faster-in-practice algorithms, and it provides inspiration for designing deterministic algorithms.

Tools from last time

An **indicator** random variable X has 2 possible outcomes: 0 and 1.

Expected value of an indicator r.v.: $\mathbb{E}[X] = \Pr[X=1]$.

Linearity of Expectation: For any (not necessarily independent!) random variables N_i :

$$\mathbb{E}\left[\sum_i N_i\right] = \sum_i \mathbb{E}[N_i].$$

Markov's Inequality: If X is a non-negative random variable and $a > 0$, then $\Pr[X \geq a] \leq \mathbb{E}[X]/a$.


QuickSort

QuickSort is a commonly used randomized sorting algorithm.

1. Pick an array element as the **“pivot”**.
2. Compare pivot to each element to partition list into **two parts**: elements **less than pivot** and elements **greater than pivot**.
3. **Recurse** on both parts of list.

How do you choose the **pivot**?

We will analyze a common strategy: **choose uniformly at random!**



Quick sort with Hungarian, folk dance


153K views • 10 years ago



kamal yassin


<https://www.youtube.com/watch?v=3San3uKKHgg>

Time Analysis

- Cost of partitioning N elements: $O(N)$
- Worst case: pivot always leaves one side empty
 - $T(N) = N + T(N - 1) + T(0)$
 - $T(N) = N + T(N - 1) + c$ [since $T(0)$ is $O(1)$]
 - $T(N) \sim N^2/2 \Rightarrow O(N^2)$ [via substitution]
- Best case: pivot divides elements equally
 - $T(N) = N + T(N / 2) + T(N / 2)$
 - $T(N) = N + 2T(N / 2) = N + 2(N / 2) + 4(N / 4) + \dots + O(1)$
 - $T(N) \sim N \log N \Rightarrow O(N \log N)$ [master theorem or substitution]
- Average case: tricky  We have the background for this now
 - Between $2N \log N$ and $\sim 1.39 N \log N \Rightarrow O(N \log N)$


Quicksort

Our goal: Prove that *for any input*, the expected running time of QuickSort is $O(n \log n)$.



Analysis of Quicksort

Question: Which pairs of elements were compared during the execution of the algorithm?




Analysis of Quicksort

- * **Idea:** Fix all the randomness at the beginning. Randomly assign **priorities** to each element, pick the one with smallest priority when choosing a pivot.


[illegible]

Analysis of Quicksort

Remember, A is the *sorted* array, not the original array



* **Example:** Sorted sequence A with priorities:

						i						j	
1	2	4	5	9	11	13	15	22	24	28	31	35	39
9	7	3	10	12	2	8	14	13	1	11	4	6	5

* **Question:** Can you tell if $A[i]$ and $A[j]$ were compared during the algorithm, just by looking at the priorities?

Analysis of Quicksort

Our goal: Prove that *for any input*, the expected running time of QuickSort is $O(n \log n)$.

We'll apply the method of indicator random variables + linearity of expectation from last time

Analysis of Quicksort

Let X be the number of comparisons made by QuickSort.

Goal: calculate $E[X]$ (since running time of Quicksort = $O(X)$)

Let X_{ij} be an indicator r.v. for whether $A[i]$ and $A[j]$ are compared.


Observe that $X = \sum_{i=1}^n \sum_{j=i+1}^n X_{ij}$.

Let's use linearity of expectation to calculate $E[X]$.

(question from 2 slides ago will be useful)

Now for a randomized *data structure*...

Skip Lists



A **dictionary** is an abstract data type (ADT) that supports insert, find, and delete operations.

Simple implementations: linked list, array

no in-order traversal

complicated bookkeeping

Faster implementations: hash tables (various types), balanced binary search trees (AVL, red-black, etc.)

Skip lists are **simpler to implement** (no need for rotations, invariants, bookkeeping)

A **skip list** is like an embellished version of a **linked list**, with expected $O(\log n)$ search time (instead of $O(n)$)



Discovered by
William Pugh, 1989

Usages of Skip Lists

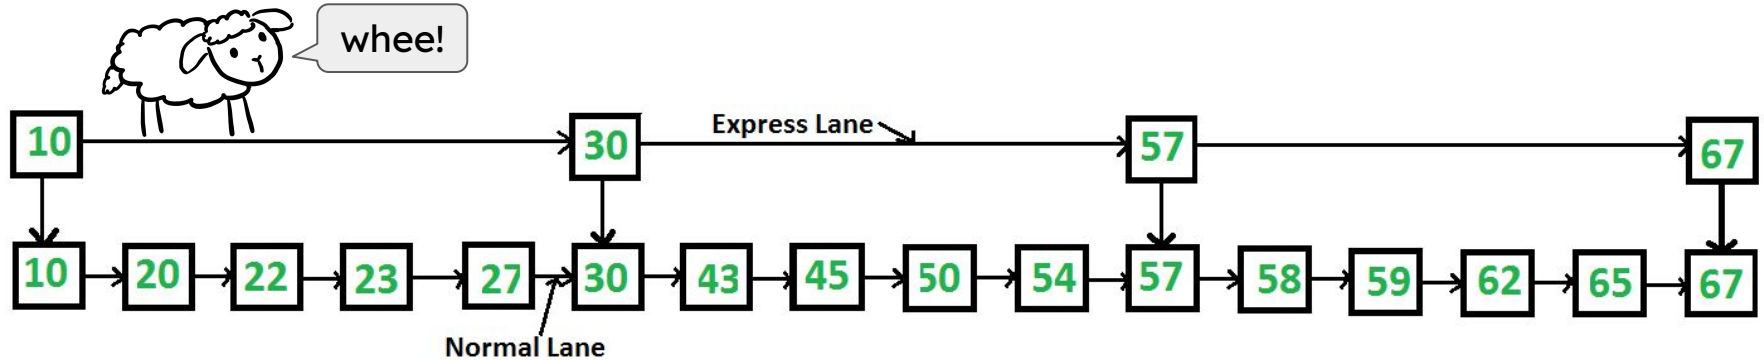
Usages [\[edit \]](#)

List of applications and frameworks that use skip lists:

- [Apache Portable Runtime](#) implements skip lists.^[9]
- [MemSQL](#) uses lock-free skip lists as its prime indexing structure for its database technology.
- [MuQSS](#), for the [Linux kernel](#), is a [cpu scheduler](#) built on skip lists.^{[10][11]}
- [Cyrus IMAP server](#) offers a "skiplist" backend DB implementation^[12]
- [Lucene](#) uses skip lists to search delta-encoded posting lists in logarithmic time.^[citation needed]
- The "QMap" key/value dictionary (up to Qt 4) template class of [Qt](#) is implemented with skip lists.^[13]
- [Redis](#), an ANSI-C open-source persistent key/value store for Posix systems, uses skip lists in its implementation of ordered sets.^[14]
- [Discord](#) uses skip lists to handle storing and updating the list of members in a [server](#).^[15]
- [RocksDB](#) uses skip lists for its default Memtable implementation.^[16]

Warm-up: 2-level Linked List

A linked list with an “express lane”

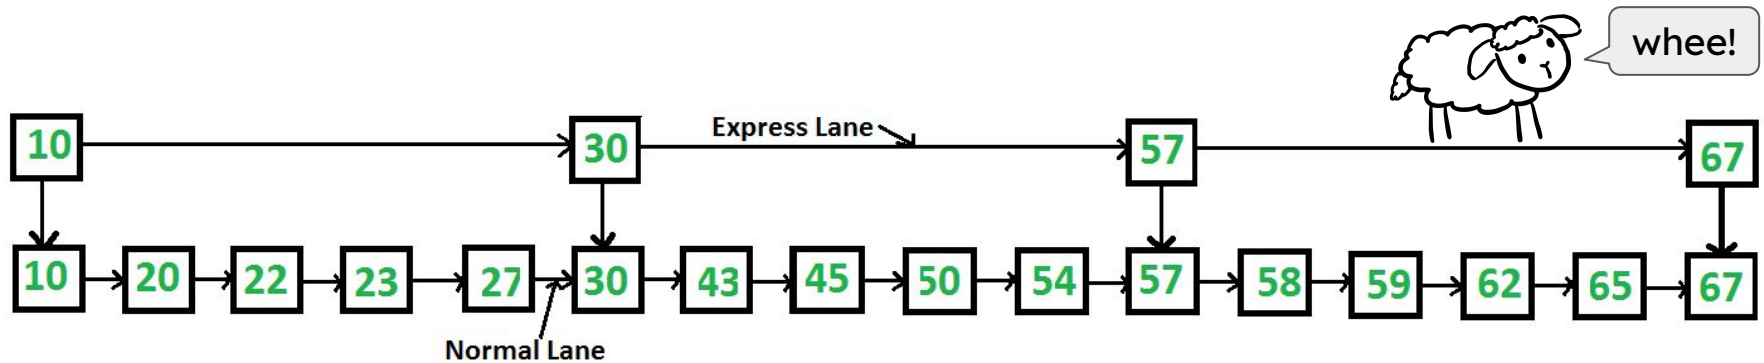


Worst-case running time for search?

How many elements should be in the express lane?

Warm-up: 2-level Linked List

A linked list with an “express lane”



Deterministically keeping express lane roughly evenly spaced amidst insertions/deletions would require some bookkeeping...

Instead let's promote elements to the express lane randomly!

With what probability should we promote each element?

Idea of a Skip List: Add more express lanes!

(Roughly $\log n$ of them)

Put every element in **level 1**

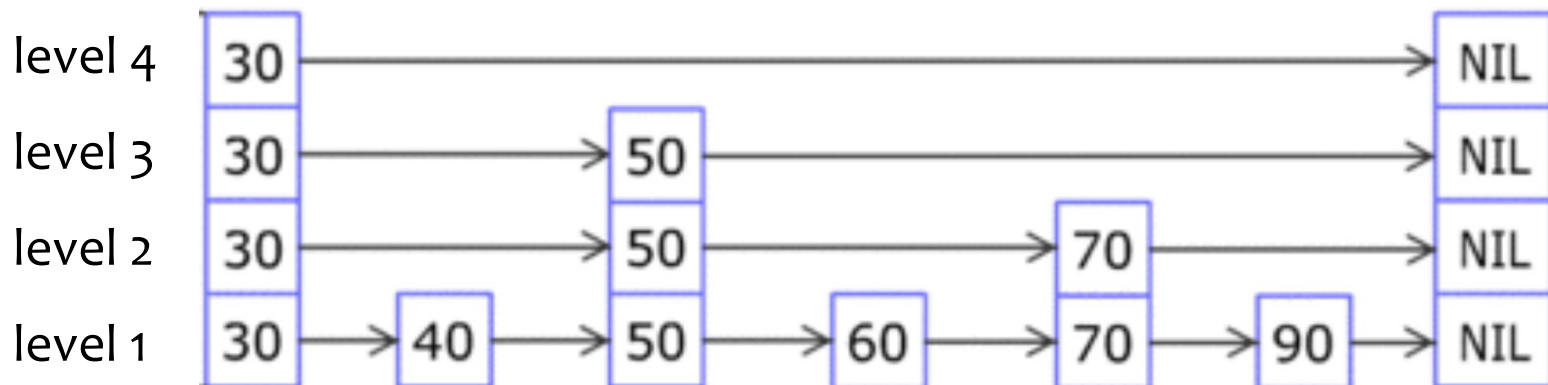
Promote about $\frac{1}{2}$ of the elements to **level 2**

Promote about $\frac{1}{2}$ of the elements in **level 2** to **level 3**

Promote about $\frac{1}{2}$ of the elements in **level 3** to **level 4**

⋮

How to accomplish this: For each element, flip a coin and keep promoting until you get tails.



Skip Lists

How do we insert an element?
(delete and find are very similar to insert)



Skip Lists

We will prove: *For any fixed sequence of n operations (insert, delete, find), the expected running time per operation is $O(\log n)$.*

Note: Our analysis will rely on the fact that the choice of operations cannot respond to the random choices of the algorithm.

(Imagine the sequence of operations is fixed ahead of time.)

Properties of Skip Lists

Goal #1: Show that the expected number of levels is $O(\log n)$.

Suppose we have a skip list containing elements x_1, x_2, \dots

Question 1: In expectation, how many elements are on level i ?

What tool should we use to answer this question?

Fix a level i . Let X be the number of elements on level i .

Properties of Skip Lists

Goal #1: Show that the expected number of levels is $O(\log n)$.

Fix a level i . Let X be the number of elements on level i .

Question 2: What is the probability level i has at least one element?

What tool should we use to answer this question?

Properties of Skip Lists

Goal #1: Show that the expected number of levels is $O(\log n)$.

Question 3: In expectation, how many levels have ≥ 1 element?

Let Y_k be an indicator that **level $\log_2 n + k$** has ≥ 1 element.

$$\begin{aligned} E[\# \text{ levels above } \log_2 n] &= E\left[\sum_{k=1}^{\infty} Y_k\right] \\ &= \sum_{k=1}^{\infty} E[Y_k] && \text{(linearity of expectation)} \\ &= \sum_{k=1}^{\infty} \Pr[Y_k=1] && \text{(expectation of indicator)} \\ &\leq \end{aligned}$$

Properties of Skip Lists

Goal #2: Show that in expectation each operation (insert, delete, find) touches a constant number of elements per level.

Combining both goals, by linearity of expectation we've proved the original goal:

Original Goal: *For any fixed sequence of n operations (insert, delete, find), the expected running time per operation is $O(\log n)$.*

Properties of Skip Lists

Goal #2: Show that in expectation each operation (insert, delete find) touches a constant number of elements per level.

Say we are inserting (or deleting or finding) y .

Fix a level i . Let Z be the number of elements touched on level i .

Let Z_j be an indicator variable for whether e_j is touched (on level i)


$$E[Z] = 2 + E[\sum_j Z_j]$$

$$= 2 + \sum_j E[Z_j] \quad (\text{linearity of expectation})$$

$$= 2 + \sum_j \Pr[Z_j=1] \quad (\text{expectation of indicator})$$

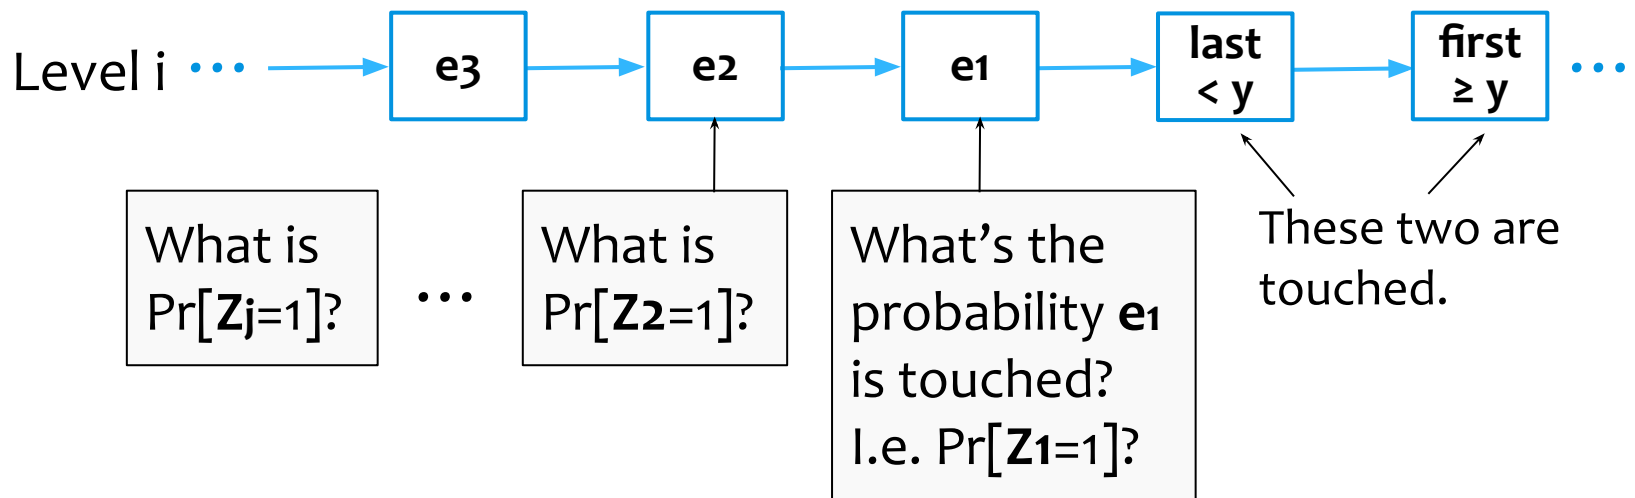
$$= 2 +$$

Claim: These two are touched. Why?



Properties of Skip Lists

Goal #2: Show that in expectation each operation (insert, delete find) touches a constant number of elements per level.



What we showed

Goal #1: Show that the expected number of levels is $O(\log n)$.

Goal #2: Show that in expectation each operation (insert, delete, find) touches a constant number of elements per level.

Original Goal: *For any fixed sequence of n operations (insert, delete, find), the expected running time per operation is $O(\log n)$.*

High-level takeaway:

- Often, deterministic data structures use a lot of bookkeeping to maintain the desired properties.
- With randomization, we stop micromanaging our data structure and use random choices that satisfy the desired properties on average.

